NetApp

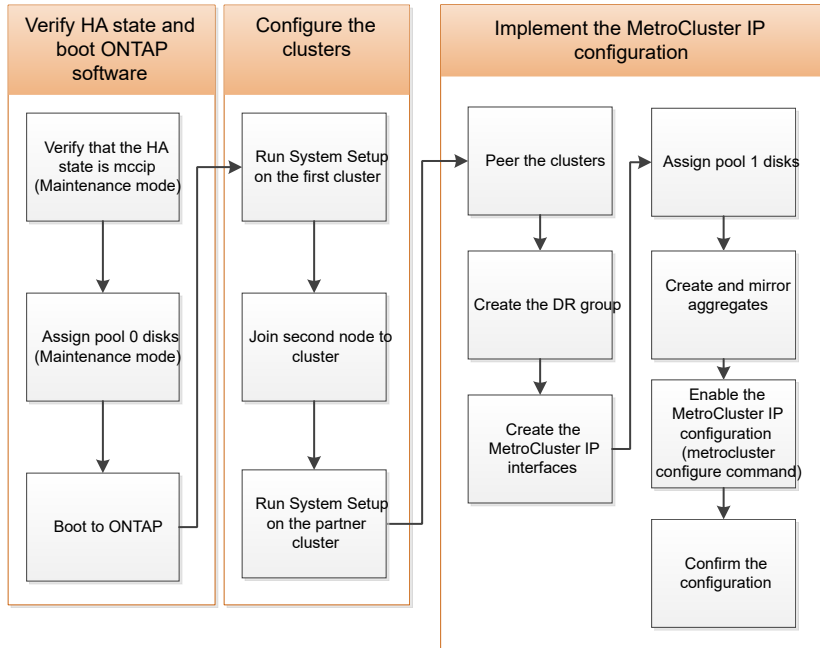# Configure the MetroCluster software in ONTAP

ONTAP MetroCluster

NetApp
April 13, 2022

Table of Contents

Configure the MetroCluster software in ONTAP

Configuring the MetroCluster software in ONTAP

You must set up each node in the MetroCluster configuration in ONTAP, including the node-level configurations and the configuration of the nodes into two sites. You must also implement the MetroCluster relationship between the two sites.

Verify HA state and boot ONTAP software	Configure the clusters	Implement the MetroCluster IP configuration	
Verify that the HA state is mccip (Maintenance mode)	Run System Setup on the first cluster	Peer the clusters	Assign pool 1 disks
Assign pool 0 disks (Maintenance mode)	Join second node to cluster	Create the DR group	Create and mirror aggregates
Boot to ONTAP	Run System Setup on the partner cluster	Create the MetroCluster IP interfaces	Enable the MetroCluster IP configuration (metrocluster configure command)
			Confirm the configuration

Handling eight-node Configurations

An eight-node configuration will consist of two DR groups. Configure the first DR group by using the tasks in this section.

Then perform the tasks in Expanding a four-node MetroCluster IP configuration to an eight-node configuration

Gathering required information

You need to gather the required IP addresses for the controller modules before you begin the configuration process.

You can use these links to download csv files and fill in the tables with your site-specific information.

MetroCluster IP setup worksheet, site_A

MetroCluster IP setup worksheet, site_B

Similarities and differences between standard cluster and MetroCluster configurations

The configuration of the nodes in each cluster in a MetroCluster configuration is similar to that of nodes in a standard cluster.

The MetroCluster configuration is built on two standard clusters. Physically, the configuration must be symmetrical, with each node having the same hardware configuration, and all of the MetroCluster components must be cabled and configured. However, the basic software configuration for nodes in a MetroCluster configuration is the same as that for nodes in a standard cluster.

Configuration step	Standard cluster configuration	MetroCluster configuration
Configure management, cluster, and data LIFs on each node.	Same in both types of clusters	
Configure the root aggregate.	Same in both types of clusters	
Set up the cluster on one node in the cluster.	Same in both types of clusters	
Join the other node to the cluster.	Same in both types of clusters	
Create a mirrored root aggregate.	Optional	Required
Peer the clusters.	Optional	Required
Enable the MetroCluster configuration.	Does not apply	Required

Verifying the ha-config state of components

In a MetroCluster IP configuration that is not preconfigured at the factory, you must verify that the ha-config state of the controller and chassis components is set to "mccip" so that they boot up properly. For systems received from the factory, this value is preconfigured and you do not need to verify it.

Before you begin
The system must be in Maintenance mode.

Steps

1. Display the HA state of the controller module and chassis:

   ```
   ha-config show
   ```

   The controller module and chassis should show the value "mccip".

2. If the displayed system state of the controller is not "mccip", set the HA state for the controller:

```
ha-config modify controller mccip
```

3. If the displayed system state of the chassis is not "mccip", set the HA state for the chassis:

```
ha-config modify chassis mccip
```

4. Repeat these steps on each node in the MetroCluster configuration.

Restoring system defaults on a controller module

Reset and restore defaults on the controller modules.

1. At the LOADER prompt, return environmental variables to their default setting: `set-defaults`

2. Boot the node to the boot menu: `boot_ontap menu`

   After you run this command, wait until the boot menu is shown.

3. Clear the node configuration:

   ◦ If you are using systems configured for ADP, select option `9a` from the boot menu, and respond `yes` when prompted.

   (i)   This process is disruptive.

   The following screen shows the boot menu prompt:

```
 Please choose one of the following:

     (1)  Normal Boot.
     (2)  Boot without /etc/rc.
     (3)  Change password.
     (4)  Clean configuration and initialize all disks.
     (5)  Maintenance mode boot.
     (6)  Update flash from backup config.
     (7)  Install new software first.
     (8)  Reboot node.
     (9)  Configure Advanced Drive Partitioning.
     Selection (1-9)?  9a
########## WARNING ##########

     This is a disruptive operation and will result in the
     loss of all filesystem data. Before proceeding further,
     make sure that:
     1) This option (9a) has been executed or will be executed
     on the HA partner node, prior to reinitializing either
     system in the HA-pair.
     2) The HA partner node is currently in a halted state or
     at the LOADER prompt.


     Do you still want to continue (yes/no)? yes
```

◦ If your system is not configured for ADP, type `wipeconfig` at the boot menu prompt, and then press Enter.

  The following screen shows the boot menu prompt:

```
Please choose one of the following:

    (1)  Normal Boot.
    (2)  Boot without /etc/rc.
    (3)  Change password.
    (4)  Clean configuration and initialize all disks.
    (5)  Maintenance mode boot.
    (6)  Update flash from backup config.
    (7)  Install new software first.
    (8)  Reboot node.
    (9)  Configure Advanced Drive Partitioning.
    Selection (1-9)?  wipeconfig
This option deletes critical system configuration, including cluster
membership.
Warning: do not run this option on a HA node that has been taken
over.
Are you sure you want to continue?: yes
Rebooting to finish wipeconfig request.
```

Manually assigning drives to pool 0

If you did not receive the systems pre-configured from the factory, you might have to manually assign the pool 0 drives. Depending on the platform model and whether the system is using ADP, you must manually assign drives to pool 0 for each node in the MetroCluster IP configuration. The procedure you use depends on the version of ONTAP you are using.

- Manually assigning drives for pool 0 (ONTAP 9.4 and later)
- Manually assigning drives for pool 0 (ONTAP 9.3)

Manually assigning drives for pool 0 (ONTAP 9.4 and later)

If the system has not been pre-configured at the factory and does not meet the requirements for automatic drive assignment, you must manually assign the pool 0 drives.

About this task

This procedure applies to configurations running ONTAP 9.4 or later.

To determine if your system requires manual disk assignment, you should review Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later.

You perform these steps in Maintenance mode. The procedure must be performed on each node in the configuration.

Examples in this section are based on the following assumptions:

- node_A_1 and node_A_2 own drives on:
  - site_A-shelf_1 (local)

- ◦ site_B-shelf_2 (remote)
- • node_B_1 and node_B_2 own drives on:
  - ◦ site_B-shelf_1 (local)
  - ◦ site_A-shelf_2 (remote)

Steps

1. Display the boot menu:

   ```
   boot_ontap menu
   ```

2. Select option "9a".

   The following screen shows the boot menu prompt:

   ```
   Please choose one of the following:

       (1)  Normal Boot.
       (2)  Boot without /etc/rc.
       (3)  Change password.
       (4)  Clean configuration and initialize all disks.
       (5)  Maintenance mode boot.
       (6)  Update flash from backup config.
       (7)  Install new software first.
       (8)  Reboot node.
       (9)  Configure Advanced Drive Partitioning.
       Selection (1-9)?  9a
   ########## WARNING ##########

       This is a disruptive operation and will result in the
       loss of all filesystem data. Before proceeding further,
       make sure that:
       1) This option (9a) has been executed or will be executed
       on the HA partner node (and DR/DR-AUX partner nodes if
       applicable), prior to reinitializing any system in the
       HA-pair (or MetroCluster setup).
       2) The HA partner node (and DR/DR-AUX partner nodes if
       applicable) is currently waiting at the boot menu.

       Do you still want to continue (yes/no)? yes
   ```

3. When the node restarts, press Ctrl-C when prompted to display the boot menu, and then select the option for **Maintenance mode boot**.

4. In Maintenance mode, manually assign drives for the local aggregates on the node:

   ```
   disk assign disk-id -p 0 -s local-node-sysid
   ```

The drives should be assigned symmetrically, so each node has an equal number of drives. The following steps are for a configuration with two storage shelves at each site.

    a. When configuring node_A_1, manually assign drives from slot 0 to 11 to pool0 of node A1 from site_A-shelf_1.

    b. When configuring node_A_2, manually assign drives from slot 12 to 23 to pool0 of node A2 from site_A-shelf_1.

    c. When configuring node_B_1, manually assign drives from slot 0 to 11 to pool0 of node B1 from site_B-shelf_1.

    d. When configuring node_B_2, manually assign drives from slot 12 to 23 to pool0 of node B2 from site_B-shelf_1.

5. Exit Maintenance mode:

```
halt
```

6. Display the boot menu:

```
boot_ontap menu
```

7. Select option "4" from the boot menu and let the system boot.

8. Repeat these steps on the other nodes in the MetroCluster IP configuration.

9. Proceed to Setting up ONTAP.

Manually assigning drives for pool 0 (ONTAP 9.3)

If you have at least two disk shelves for each node, you use ONTAP's auto-assignment functionality to automatically assign the local (pool 0) disks.

About this task

While the node is in Maintenance mode, you must first assign a single disk on the appropriate shelves to pool 0. ONTAP then automatically assign the rest of the disks on the shelf to the same pool. This task is not required on systems received from the factory, which have pool 0 to contain the pre-configured root aggregate.

This procedure applies to configurations running ONTAP 9.3.

This procedure is not required if you received your MetroCluster configuration from the factory. Nodes from the factory are configured with pool 0 disks and root aggregates.

This procedure can be used only if you have at least two disk shelves for each node, which allows shelf-level autoassignment of disks. If you cannot use shelf-level autoassignment, you must manually assign your local disks so that each node has a local pool of disks (pool 0).

These steps must be performed in Maintenance mode.

Examples in this section assume the following disk shelves:

- node_A_1 owns disks on:
  - site_A-shelf_1 (local)
  - site_B-shelf_2 (remote)
- node_A_2 is connected to:

- ◦ site_A-shelf_3 (local)
- ◦ site_B-shelf_4 (remote)
- • node_B_1 is connected to:
  - ◦ site_B-shelf_1 (local)
  - ◦ site_A-shelf_2 (remote)
- • node_B_2 is connected to:
  - ◦ site_B-shelf_3 (local)
  - ◦ site_A-shelf_4 (remote)

Steps

1. Manually assign a single disk for root aggregate on each node:

   ```
   disk assign disk-id -p 0 -s local-node-sysid
   ```

   The manual assignment of these disks allows the ONTAP autoassignment feature to assign the rest of the disks on each shelf.

   a. On node_A_1, manually assign one disk from local site_A-shelf_1 to pool 0.

   b. On node_A_2, manually assign one disk from local site_A-shelf_3 to pool 0.

   c. On node_B_1, manually assign one disk from local site_B-shelf_1 to pool 0.

   d. On node_B_2, manually assign one disk from local site_B-shelf_3 to pool 0.

2. Boot each node at site A, using option "4" on the boot menu:

   You should complete this step on a node before proceeding to the next node.

   a. Exit Maintenance mode:

   ```
   halt
   ```

   b. Display the boot menu:

   ```
   boot_ontap menu
   ```

   c. Select option "4" from the boot menu and proceed.

3. Boot each node at site B, using option "4" on the boot menu:

   You should complete this step on a node before proceeding to the next node.

   a. Exit Maintenance mode:

   ```
   halt
   ```

   b. Display the boot menu:

   ```
   boot_ontap menu
   ```

   c. Select option "4" from the boot menu and proceed.

Setting up ONTAP

After you boot each node, you are prompted to perform basic node and cluster configuration. After configuring the cluster, you return to the ONTAP CLI to create aggregates and create the MetroCluster configuration.

Before you begin

- You must have cabled the MetroCluster configuration.
- You must not have configured the Service Processor.

If you need to netboot the new controllers, see Netbooting the new controller modules.

About this task

This task must be performed on both clusters in the MetroCluster configuration.

Steps

1. Power up each node at the local site if you have not already done so and let them all boot completely.

   If the system is in Maintenance mode, you need to issue the halt command to exit Maintenance mode, and then issue the `boot_ontap` command to boot the system and get to cluster setup.

2. On the first node in each cluster, proceed through the prompts to configure the cluster

   a. Enable the AutoSupport tool by following the directions provided by the system.

      The output should be similar to the following:

```
Welcome to the cluster setup wizard.

    You can enter the following commands at any time:
    "help" or "?" - if you want to have a question clarified,
    "back" - if you want to change previously answered questions, and
    "exit" or "quit" - if you want to quit the cluster setup wizard.
    Any changes you made before quitting will be saved.

    You can return to cluster setup at any time by typing "cluster
setup".
    To accept a default or omit a question, do not enter a value.

    This system will send event messages and periodic reports to
NetApp Technical
    Support. To disable this feature, enter
    autosupport modify -support disable
    within 24 hours.

    Enabling AutoSupport can significantly speed problem
determination and
    resolution should a problem occur on your system.
    For further information on AutoSupport, see:
    http://support.netapp.com/autosupport/

    Type yes to confirm and continue {yes}: yes

.
.
.
```

b. Configure the node management interface by responding to the prompts.

The prompts are similar to the following:

```
Enter the node management interface port [e0M]:
Enter the node management interface IP address: 172.17.8.229
Enter the node management interface netmask: 255.255.254.0
Enter the node management interface default gateway: 172.17.8.1
A node management interface on port e0M with IP address 172.17.8.229
has been created.
```

c. Create the cluster by responding to the prompts.

The prompts are similar to the following:

```
Do you want to create a new cluster or join an existing cluster?
{create, join}:
create


Do you intend for this node to be used as a single node cluster?
{yes, no} [no]:
no

Existing cluster interface configuration found:

Port MTU IP Netmask
e0a 1500 169.254.18.124 255.255.0.0
e1a 1500 169.254.184.44 255.255.0.0

Do you want to use this configuration? {yes, no} [yes]: no

System Defaults:
Private cluster network ports [e0a,e1a].
Cluster port MTU values will be set to 9000.
Cluster interface IP addresses will be automatically generated.

Do you want to use these defaults? {yes, no} [yes]: no

Enter the cluster administrator's (username "admin") password:

Retype the password:


Step 1 of 5: Create a Cluster
You can type "back", "exit", or "help" at any question.

List the private cluster network ports [e0a,e1a]:
Enter the cluster ports' MTU size [9000]:
Enter the cluster network netmask [255.255.0.0]: 255.255.254.0
Enter the cluster interface IP address for port e0a: 172.17.10.228
Enter the cluster interface IP address for port e1a: 172.17.10.229
Enter the cluster name: cluster_A

Creating cluster cluster_A

Starting cluster support services ...

Cluster cluster_A has been created.
```

d. Add licenses, set up a Cluster Administration SVM, and enter DNS information by responding to the prompts.

The prompts are similar to the following:

```
Step 2 of 5: Add Feature License Keys
You can type "back", "exit", or "help" at any question.


Enter an additional license key []:



Step 3 of 5: Set Up a Vserver for Cluster Administration
You can type "back", "exit", or "help" at any question.



Enter the cluster management interface port [e3a]:
Enter the cluster management interface IP address: 172.17.12.153
Enter the cluster management interface netmask: 255.255.252.0
Enter the cluster management interface default gateway: 172.17.12.1

A cluster management interface on port e3a with IP address
172.17.12.153 has been created. You can use this address to connect
to and manage the cluster.

Enter the DNS domain names: lab.netapp.com
Enter the name server IP addresses: 172.19.2.30
DNS lookup for the admin Vserver will use the lab.netapp.com domain.

Step 4 of 5: Configure Storage Failover (SFO)
You can type "back", "exit", or "help" at any question.



SFO will be enabled when the partner joins the cluster.


Step 5 of 5: Set Up the Node
You can type "back", "exit", or "help" at any question.

Where is the controller located []: svl
```

e. Enable storage failover and set up the node by responding to the prompts.

The prompts are similar to the following:

```
Step 4 of 5: Configure Storage Failover (SFO)
You can type "back", "exit", or "help" at any question.


SFO will be enabled when the partner joins the cluster.


Step 5 of 5: Set Up the Node
You can type "back", "exit", or "help" at any question.

Where is the controller located []: site_A
```

    f. Complete the configuration of the node, but do not create data aggregates.

       You can use ONTAP System Manager, pointing your web browser to the cluster management IP address (https://172.17.12.153).

       Cluster management using System Manager (Versions 9.0 to 9.6)

       ONTAP System Manager (Versions 9.7 and later)

3. Boot the next controller and join it to the cluster, following the prompts.

4. Confirm that nodes are configured in high-availability mode:

```
storage failover show -fields mode
```

If not, you must configure HA mode on each node, and then reboot the nodes:

```
storage failover modify -mode ha -node localhost
```

This command configures high-availability mode but does not enable storage failover. Storage failover is automatically enabled when you configure the MetroCluster configuration later in the process.

5. Confirm that you have four ports configured as cluster interconnects:

```
network port show
```

The MetroCluster IP interfaces are not configured at this time and do not appear in the command output.

The following example shows two cluster ports on node_A_1:

```
cluster_A::*> network port show -role cluster



Node: node_A_1


Ignore
```

```
                                                      Speed(Mbps) Health
Health

Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status

--------- ------------ ---------------- ---- ---- ----------- --------
------

e4a       Cluster      Cluster          up   9000 auto/40000 healthy
false

e4e       Cluster      Cluster          up   9000 auto/40000 healthy
false


Node: node_A_2


Ignore

                                                      Speed(Mbps) Health
Health

Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status

--------- ------------ ---------------- ---- ---- ----------- --------
------

e4a       Cluster      Cluster          up   9000 auto/40000 healthy
false

e4e       Cluster      Cluster          up   9000 auto/40000 healthy
false


4 entries were displayed.
```

6. Repeat these steps on the partner cluster.

Return to the ONTAP command-line interface and complete the MetroCluster configuration by performing the tasks that follow.

Configuring the clusters into a MetroCluster configuration

You must peer the clusters, mirror the root aggregates, create a mirrored data aggregate, and then issue the command to implement the MetroCluster operations.

Disabling automatic drive assignment (if doing manual assignment in ONTAP 9.4)

In ONTAP 9.4, if your MetroCluster IP configuration has fewer than four external storage shelves per site, you must disable automatic drive assignment on all nodes and manually assign drives.

About this task

This task is not required in ONTAP 9.5 and later.

This task does not apply to an AFF A800 system with an internal shelf and no external shelves.

Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later

Steps

1. Disable automatic drive assignment:

   ```
   storage disk option modify -node node_name -autoassign off
   ```

   You need to issue this command on all nodes in the MetroCluster IP configuration.

Verifying drive assignment of pool 0 drives

You must verify that the remote drives are visible to the nodes and have been assigned correctly.

About this task

Automatic assignment depends on the storage system platform model and drive shelf arrangement.

Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later

Steps

1. Verify that pool 0 drives are assigned automatically:

   ```
   disk show
   ```

   The following example shows the "cluster_A" output for an AFF A800 system with no external shelves.

   One quarter (8 drives) were automatically assigned to "node_A_1" and one quarter were automatically assigned to "node_A_2". The remaining drives will be remote (pool 1) drives for "node_B_1" and "node_B_2".

   ```
   cluster_A::*> disk show
                   Usable     Disk     Container          Container
   Disk            Size       Shelf Bay Type     Type      Name
   Owner
   ---------------- ---------- ----- --- ------- ----------- ---------
   --------
   node_A_1:0n.12   1.75TB     0     12  SSD-NVM shared      aggr0
   ```

```
node_A_1
node_A_1:0n.13    1.75TB       0       13    SSD-NVM shared       aggr0
node_A_1
node_A_1:0n.14    1.75TB       0       14    SSD-NVM shared       aggr0
node_A_1
node_A_1:0n.15    1.75TB       0       15    SSD-NVM shared       aggr0
node_A_1
node_A_1:0n.16    1.75TB       0       16    SSD-NVM shared       aggr0
node_A_1
node_A_1:0n.17    1.75TB       0       17    SSD-NVM shared       aggr0
node_A_1
node_A_1:0n.18    1.75TB       0       18    SSD-NVM shared       aggr0
node_A_1
node_A_1:0n.19    1.75TB       0       19    SSD-NVM shared       -
node_A_1
node_A_2:0n.0     1.75TB       0       0     SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.1     1.75TB       0       1     SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.2     1.75TB       0       2     SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.3     1.75TB       0       3     SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.4     1.75TB       0       4     SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.5     1.75TB       0       5     SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.6     1.75TB       0       6     SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.7     1.75TB       0       7     SSD-NVM shared       -
node_A_2
node_A_2:0n.24    -            0       24    SSD-NVM unassigned   -          -
node_A_2:0n.25    -            0       25    SSD-NVM unassigned   -          -
node_A_2:0n.26    -            0       26    SSD-NVM unassigned   -          -
node_A_2:0n.27    -            0       27    SSD-NVM unassigned   -          -
node_A_2:0n.28    -            0       28    SSD-NVM unassigned   -          -
node_A_2:0n.29    -            0       29    SSD-NVM unassigned   -          -
node_A_2:0n.30    -            0       30    SSD-NVM unassigned   -          -
node_A_2:0n.31    -            0       31    SSD-NVM unassigned   -          -
node_A_2:0n.36    -            0       36    SSD-NVM unassigned   -          -
node_A_2:0n.37    -            0       37    SSD-NVM unassigned   -          -
node_A_2:0n.38    -            0       38    SSD-NVM unassigned   -          -
node_A_2:0n.39    -            0       39    SSD-NVM unassigned   -          -
node_A_2:0n.40    -            0       40    SSD-NVM unassigned   -          -
node_A_2:0n.41    -            0       41    SSD-NVM unassigned   -          -
node_A_2:0n.42    -            0       42    SSD-NVM unassigned   -          -
```

```
node_A_2:0n.43    -            0      43  SSD-NVM unassigned  -            -
32 entries were displayed.
```

The following example shows the "cluster_B" output:

```
cluster_B::> disk show
                  Usable    Disk               Container  Container
Disk              Size      Shelf Bay Type      Type       Name
Owner
--------------- ---------- ----- --- ------- ----------- ---------
--------

Info: This cluster has partitioned disks. To get a complete list of
spare disk
capacity use "storage aggregate show-spare-disks".
node_B_1:0n.12    1.75TB     0      12  SSD-NVM shared       aggr0
node_B_1
node_B_1:0n.13    1.75TB     0      13  SSD-NVM shared       aggr0
node_B_1
node_B_1:0n.14    1.75TB     0      14  SSD-NVM shared       aggr0
node_B_1
node_B_1:0n.15    1.75TB     0      15  SSD-NVM shared       aggr0
node_B_1
node_B_1:0n.16    1.75TB     0      16  SSD-NVM shared       aggr0
node_B_1
node_B_1:0n.17    1.75TB     0      17  SSD-NVM shared       aggr0
node_B_1
node_B_1:0n.18    1.75TB     0      18  SSD-NVM shared       aggr0
node_B_1
node_B_1:0n.19    1.75TB     0      19  SSD-NVM shared       -
node_B_1
node_B_2:0n.0     1.75TB     0      0   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.1     1.75TB     0      1   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.2     1.75TB     0      2   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.3     1.75TB     0      3   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.4     1.75TB     0      4   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.5     1.75TB     0      5   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.6     1.75TB     0      6   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
```

```
    node_B_2:0n.7     1.75TB     0      7    SSD-NVM shared      -
    node_B_2
    node_B_2:0n.24    -          0      24   SSD-NVM unassigned  -        -
    node_B_2:0n.25    -          0      25   SSD-NVM unassigned  -        -
    node_B_2:0n.26    -          0      26   SSD-NVM unassigned  -        -
    node_B_2:0n.27    -          0      27   SSD-NVM unassigned  -        -
    node_B_2:0n.28    -          0      28   SSD-NVM unassigned  -        -
    node_B_2:0n.29    -          0      29   SSD-NVM unassigned  -        -
    node_B_2:0n.30    -          0      30   SSD-NVM unassigned  -        -
    node_B_2:0n.31    -          0      31   SSD-NVM unassigned  -        -
    node_B_2:0n.36    -          0      36   SSD-NVM unassigned  -        -
    node_B_2:0n.37    -          0      37   SSD-NVM unassigned  -        -
    node_B_2:0n.38    -          0      38   SSD-NVM unassigned  -        -
    node_B_2:0n.39    -          0      39   SSD-NVM unassigned  -        -
    node_B_2:0n.40    -          0      40   SSD-NVM unassigned  -        -
    node_B_2:0n.41    -          0      41   SSD-NVM unassigned  -        -
    node_B_2:0n.42    -          0      42   SSD-NVM unassigned  -        -
    node_B_2:0n.43    -          0      43   SSD-NVM unassigned  -        -
    32 entries were displayed.

    cluster_B::>
```

Peering the clusters

The clusters in the MetroCluster configuration must be in a peer relationship so that they can communicate with each other and perform the data mirroring essential to MetroCluster disaster recovery.

Related information

Cluster and SVM peering express configuration

Considerations when using dedicated ports

Considerations when sharing data ports

Configuring intercluster LIFs for cluster peering

You must create intercluster LIFs on ports used for communication between the MetroCluster partner clusters. You can use dedicated ports or ports that also have data traffic.

Configuring intercluster LIFs on dedicated ports

You can configure intercluster LIFs on dedicated ports. Doing so typically increases the available bandwidth for replication traffic.

Steps

1. List the ports in the cluster:

    ```
    network port show
    ```

For complete command syntax, see the man page.

The following example shows the network ports in cluster01:

```
cluster01::> network port show
                                                          Speed
 (Mbps)
Node    Port      IPspace        Broadcast Domain Link   MTU    Admin/Oper
------  --------- ------------   ---------------- -----  -------
------------
cluster01-01
        e0a       Cluster        Cluster              up     1500   auto/1000
        e0b       Cluster        Cluster              up     1500   auto/1000
        e0c       Default        Default              up     1500   auto/1000
        e0d       Default        Default              up     1500   auto/1000
        e0e       Default        Default              up     1500   auto/1000
        e0f       Default        Default              up     1500   auto/1000
cluster01-02
        e0a       Cluster        Cluster              up     1500   auto/1000
        e0b       Cluster        Cluster              up     1500   auto/1000
        e0c       Default        Default              up     1500   auto/1000
        e0d       Default        Default              up     1500   auto/1000
        e0e       Default        Default              up     1500   auto/1000
        e0f       Default        Default              up     1500   auto/1000
```

2. Determine which ports are available to dedicate to intercluster communication:

```
network interface show -fields home-port,curr-port
```

For complete command syntax, see the man page.

The following example shows that ports "e0e" and "e0f" have not been assigned LIFs:

```
cluster01::> network interface show -fields home-port,curr-port
vserver lif                        home-port curr-port
------- -------------------- --------- ---------
Cluster cluster01-01_clus1    e0a       e0a
Cluster cluster01-01_clus2    e0b       e0b
Cluster cluster01-02_clus1    e0a       e0a
Cluster cluster01-02_clus2    e0b       e0b
cluster01
        cluster_mgmt          e0c       e0c
cluster01
        cluster01-01_mgmt1    e0c       e0c
cluster01
        cluster01-02_mgmt1    e0c       e0c
```

3. Create a failover group for the dedicated ports:

```
network interface failover-groups create -vserver system_SVM -failover-group
failover_group -targets physical_or_logical_ports
```

The following example assigns ports "e0e" and "e0f" to the failover group "intercluster01" on the system "SVMcluster01":

```
cluster01::> network interface failover-groups create -vserver cluster01
-failover-group
intercluster01 -targets
cluster01-01:e0e,cluster01-01:e0f,cluster01-02:e0e,cluster01-02:e0f
```

4. Verify that the failover group was created:

```
network interface failover-groups show
```

For complete command syntax, see the man page.

```
cluster01::> network interface failover-groups show
                                    Failover
Vserver          Group            Targets
---------------- ----------------
-----------------------------------------------
Cluster
                 Cluster
                                  cluster01-01:e0a, cluster01-01:e0b,
                                  cluster01-02:e0a, cluster01-02:e0b
cluster01
                 Default
                                  cluster01-01:e0c, cluster01-01:e0d,
                                  cluster01-02:e0c, cluster01-02:e0d,
                                  cluster01-01:e0e, cluster01-01:e0f
                                  cluster01-02:e0e, cluster01-02:e0f
                 intercluster01
                                  cluster01-01:e0e, cluster01-01:e0f
                                  cluster01-02:e0e, cluster01-02:e0f
```

5. Create intercluster LIFs on the system SVM and assign them to the failover group.

ONTAP version	Command

9.6 and later	`network interface create -vserver system_SVM -lif LIF_name -service -policy default-intercluster -home -node node -home-port port -address port_IP -netmask netmask -failover -group failover_group`
9.5 and earlier	`network interface create -vserver system_SVM -lif LIF_name -role intercluster -home-node node -home -port port -address port_IP -netmask netmask -failover-group failover_group`

For complete command syntax, see the man page.

The following example creates intercluster LIFs "cluster01_icl01" and "cluster01_icl02" in the failover group "intercluster01":

```
cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl01 -service-
policy default-intercluster -home-node cluster01-01 -home-port e0e
-address 192.168.1.201
-netmask 255.255.255.0 -failover-group intercluster01

cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl02 -service-
policy default-intercluster -home-node cluster01-02 -home-port e0e
-address 192.168.1.202
-netmask 255.255.255.0 -failover-group intercluster01
```

6. Verify that the intercluster LIFs were created:

In ONTAP 9.6 and later:
`network interface show -service-policy default-intercluster`
In ONTAP 9.5 and earlier:
`network interface show -role intercluster`

For complete command syntax, see the man page.

```
cluster01::> network interface show -service-policy default-intercluster
            Logical     Status     Network                 Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask            Node          Port
Home
----------- ---------- ---------- ------------------ -------------
------- ----
cluster01
            cluster01_icl01
                        up/up      192.168.1.201/24   cluster01-01  e0e
true
            cluster01_icl02
                        up/up      192.168.1.202/24   cluster01-02  e0f
true
```

7. Verify that the intercluster LIFs are redundant:

In ONTAP 9.6 and later:
`network interface show -service-policy default-intercluster -failover`

In ONTAP 9.5 and earlier:
`network interface show -role intercluster -failover`

For complete command syntax, see the man page.

The following example shows that the intercluster LIFs "cluster01_icl01" and "cluster01_icl02" on the "e0e" port will fail over to the "e0f" port.

```
cluster01::> network interface show -service-policy default-intercluster
-failover
          Logical          Home                      Failover        Failover
Vserver   Interface        Node:Port                 Policy          Group
-------- --------------- --------------------- --------------- --------
cluster01
        cluster01_icl01 cluster01-01:e0e    local-only
intercluster01
                            Failover Targets:  cluster01-01:e0e,
                                               cluster01-01:e0f
        cluster01_icl02 cluster01-02:e0e    local-only
intercluster01
                            Failover Targets:  cluster01-02:e0e,
                                               cluster01-02:e0f
```

Configuring intercluster LIFs on shared data ports

You can configure intercluster LIFs on ports shared with the data network. Doing so reduces the number of ports you need for intercluster networking.

1. List the ports in the cluster:

   `network port show`

   For complete command syntax, see the man page.

   The following example shows the network ports in cluster01:

   ```
   cluster01::> network port show
                                                           Speed
   (Mbps)
   Node    Port      IPspace       Broadcast Domain Link    MTU    Admin/Oper
   ------  --------  ------------  ---------------- -----  -------
   -----------
   cluster01-01
           e0a       Cluster       Cluster          up     1500    auto/1000
           e0b       Cluster       Cluster          up     1500    auto/1000
           e0c       Default       Default          up     1500    auto/1000
           e0d       Default       Default          up     1500    auto/1000
   cluster01-02
           e0a       Cluster       Cluster          up     1500    auto/1000
           e0b       Cluster       Cluster          up     1500    auto/1000
           e0c       Default       Default          up     1500    auto/1000
           e0d       Default       Default          up     1500    auto/1000
   ```

2. Create intercluster LIFs on the system SVM:

In ONTAP 9.6 and later:
`network interface create -vserver system_SVM -lif LIF_name -service-policy default-intercluster -home-node node -home-port port -address port_IP -netmask netmask`
In ONTAP 9.5 and earlier:
`network interface create -vserver system_SVM -lif LIF_name -role intercluster -home-node node -home-port port -address port_IP -netmask netmask`

   For complete command syntax, see the man page.

   The following example creates intercluster LIFs "cluster01_icl01" and "cluster01_icl02":

```
cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl01 -service-
policy default-intercluster -home-node cluster01-01 -home-port e0c
-address 192.168.1.201
-netmask 255.255.255.0

cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl02 -service-
policy default-intercluster -home-node cluster01-02 -home-port e0c
-address 192.168.1.202
-netmask 255.255.255.0
```

3. Verify that the intercluster LIFs were created:

In ONTAP 9.6 and later:
`network interface show -service-policy default-intercluster`

In ONTAP 9.5 and earlier:
`network interface show -role intercluster`

For complete command syntax, see the man page.

```
cluster01::> network interface show -service-policy default-intercluster
            Logical    Status     Network                Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask           Node          Port
Home
----------- ---------- ---------- ------------------ -------------
------- ----
cluster01
            cluster01_icl01
                       up/up      192.168.1.201/24   cluster01-01  e0c
    true
            cluster01_icl02
                       up/up      192.168.1.202/24   cluster01-02  e0c
    true
```

4. Verify that the intercluster LIFs are redundant:

In ONTAP 9.6 and later:
`network interface show –service-policy default-intercluster -failover`

```
network interface show -role intercluster -failover
```

For complete command syntax, see the man page.

The following example shows that the intercluster LIFs "cluster01_icl01" and "cluster01_icl02" on the "e0c" port will fail over to the "e0d" port.

```
cluster01::> network interface show -service-policy default-intercluster
-failover
         Logical          Home                         Failover        Failover
Vserver  Interface        Node:Port                    Policy          Group
-------- ---------------  --------------------- --------------- --------
cluster01
         cluster01_icl01 cluster01-01:e0c   local-only
192.168.1.201/24
                              Failover Targets: cluster01-01:e0c,
                                                cluster01-01:e0d
         cluster01_icl02 cluster01-02:e0c   local-only
192.168.1.201/24
                              Failover Targets: cluster01-02:e0c,
                                                cluster01-02:e0d
```

Related information

Considerations when sharing data ports

Creating a cluster peer relationship

You can use the cluster peer create command to create a peer relationship between a local and remote cluster. After the peer relationship has been created, you can run cluster peer create on the remote cluster to authenticate it to the local cluster.

Before you begin

- You must have created intercluster LIFs on every node in the clusters that are being peered.

- The clusters must be running ONTAP 9.3 or later.

Steps

1. On the destination cluster, create a peer relationship with the source cluster:

   ```
   cluster peer create -generate-passphrase -offer-expiration MM/DD/YYYY
   HH:MM:SS|1…7days|1…168hours -peer-addrs peer_LIF_IPs -ipspace ipspace
   ```

   If you specify both -generate-passphrase and -peer-addrs, only the cluster whose intercluster LIFs are specified in -peer-addrs can use the generated password.

   You can ignore the -ipspace option if you are not using a custom IPspace. For complete command

syntax, see the man page.

The following example creates a cluster peer relationship on an unspecified remote cluster:

```
cluster02::> cluster peer create -generate-passphrase -offer-expiration
2days

                       Passphrase: UCa+6lRVICXeL/gq1WrK7ShR
                  Expiration Time: 6/7/2017 08:16:10 EST
   Initial Allowed Vserver Peers: -
             Intercluster LIF IP: 192.140.112.101
               Peer Cluster Name: Clus_7ShR (temporary generated)

Warning: make a note of the passphrase - it cannot be displayed again.
```

2. On source cluster, authenticate the source cluster to the destination cluster:

```
cluster peer create -peer-addrs peer_LIF_IPs -ipspace ipspace
```

For complete command syntax, see the man page.

The following example authenticates the local cluster to the remote cluster at intercluster LIF IP addresses "192.140.112.101" and "192.140.112.102":

```
cluster01::> cluster peer create -peer-addrs
192.140.112.101,192.140.112.102

Notice: Use a generated passphrase or choose a passphrase of 8 or more
characters.
        To ensure the authenticity of the peering relationship, use a
phrase or sequence of characters that would be hard to guess.

Enter the passphrase:
Confirm the passphrase:

Clusters cluster02 and cluster01 are peered.
```

Enter the passphrase for the peer relationship when prompted.

3. Verify that the cluster peer relationship was created:

```
cluster peer show -instance
```

```
cluster01::> cluster peer show -instance

                           Peer Cluster Name: cluster02
            Remote Intercluster Addresses: 192.140.112.101,
192.140.112.102
              Availability of the Remote Cluster: Available
                         Remote Cluster Name: cluster2
                         Active IP Addresses: 192.140.112.101,
192.140.112.102
                        Cluster Serial Number: 1-80-123456
                Address Family of Relationship: ipv4
          Authentication Status Administrative: no-authentication
            Authentication Status Operational: absent
                           Last Update Time: 02/05 21:05:41
                   IPspace for the Relationship: Default
```

4. Check the connectivity and status of the nodes in the peer relationship:

`cluster peer health show`

```
cluster01::> cluster peer health show
Node         cluster-Name                Node-Name
             Ping-Status                 RDB-Health Cluster-Health  Avail…
---------- -------------------------- ---------   ---------------
--------
cluster01-01
           cluster02                   cluster02-01
             Data: interface_reachable
             ICMP: interface_reachable true        true            true
                                       cluster02-02
             Data: interface_reachable
             ICMP: interface_reachable true        true            true
cluster01-02
           cluster02                   cluster02-01
             Data: interface_reachable
             ICMP: interface_reachable true        true            true
                                       cluster02-02
             Data: interface_reachable
             ICMP: interface_reachable true        true            true
```

Creating the DR group

You must create the disaster recovery (DR) group relationships between the clusters.

About this task

You perform this procedure on one of the clusters in the MetroCluster configuration to create the DR relationships between the nodes in both clusters.

> ⓘ The DR relationships cannot be changed after the DR groups are created.

Steps

1. Verify that the nodes are ready for creation of the DR group by entering the following command on each:

   `metrocluster configuration-settings show-status`

   The command output should show that the nodes are ready:

   ```
   cluster_A::> metrocluster configuration-settings show-status
   Cluster                      Node          Configuration Settings Status
   -------------------------    ------------- 
   ------------------------------
   cluster_A                    node_A_1      ready for DR group create
                                node_A_2      ready for DR group create
   2 entries were displayed.
   ```

   ```
   cluster_B::> metrocluster configuration-settings show-status
   Cluster                      Node          Configuration Settings Status
   -------------------------    ------------- 
   ------------------------------
   cluster_B                    node_B_1      ready for DR group create
                                node_B_2      ready for DR group create
   2 entries were displayed.
   ```

2. Create the DR group:

   `metrocluster configuration-settings dr-group create -partner-cluster partner-cluster-name -local-node local-node-name -remote-node remote-node-name`

This command is issued only once. It does not need to be repeated on the partner cluster. In the command, you specify the name of the remote cluster and the name of one local node and one node on the partner cluster.

The two nodes you specify are configured as DR partners and the other two nodes (which are not specified in the command) are configured as the second DR pair in the DR group. These relationships cannot be changed after you enter this command.

The following command creates these DR pairs:

- node_A_1 and node_B_1
- node_A_2 and node_B_2

```
Cluster_A::> metrocluster configuration-settings dr-group create
-partner-cluster cluster_B -local-node node_A_1 -remote-node node_B_1
[Job 27] Job succeeded: DR Group Create is successful.
```

Configuring and connecting the MetroCluster IP interfaces

You must configure the MetroCluster IP interfaces that are used for replication of each node's storage and nonvolatile cache. You then establish the connections using the MetroCluster IP interfaces. This creates iSCSI connections for storage replication.

About this task

> (i) You must choose the MetroCluster IP addresses carefully because you cannot change them after initial configuration.

- You must create two interfaces for each node. The interfaces must be associated with the VLANs defined in the MetroCluster RCF file.

- You must create all MetroCluster IP interface "A" ports in the same VLAN and all MetroCluster IP interface "B" ports in the other VLAN. Refer to Considerations for MetroCluster IP configuration.

> (i)
> - Certain platforms use a VLAN for the MetroCluster IP interface. By default, each of the two ports use a different VLAN: 10 and 20. You can also specify a different (non-default) VLAN higher than 100 (between 101 and 4095) using the `-vlan-id parameter` in the `metrocluster configuration-settings interface create` command.
> - Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also specify the `-gateway` parameter when creating MetroCluster IP interfaces. Refer to Considerations for layer 3 wide-area networks.

The following platform models use VLANs and allow configuration of a non-default VLAN ID.

AFF platforms	FAS platforms

• AFF A220 • AFF A250 • AFF A400	• FAS2750 • FAS500f • FAS8300 • FAS8700

The following IP addresses and subnets are used in the examples:

Node	Interface	IP address	Subnet
node_A_1	MetroCluster IP interface 1	10.1.1.1	10.1.1/24
	MetroCluster IP interface 2	10.1.2.1	10.1.2/24
node_A_2	MetroCluster IP interface 1	10.1.1.2	10.1.1/24
	MetroCluster IP interface 2	10.1.2.2	10.1.2/24
node_B_1	MetroCluster IP interface 1	10.1.1.3	10.1.1/24
	MetroCluster IP interface 2	10.1.2.3	10.1.2/24
node_B_2	MetroCluster IP interface 1	10.1.1.4	10.1.1/24
	MetroCluster IP interface 2	10.1.2.4	10.1.2/24

The physical ports used by the MetroCluster IP interfaces depends on the platform model, as shown in the following table.

Platform model	MetroCluster IP port	Note
AFF A900	e5b	
	e7b	
AFF A800	e0b	
	e1b	

Platform model	MetroCluster IP port	Note
AFF A700 and FAS900	e5a	
	e5b	
AFF A400	e3a	
	e3b	
AFF A320	e0g	
	e0h	
AFF A300 and FAS8200	e1a	
	e1b	
AFF A220 and FAS2750	e0a	On these systems, these physical ports are also used as cluster interfaces.
	e0b	
AFF A250 and FAS500f	e0c	
	e0d	
FAS8300 and FAS8700	e0c	
	e0d	

The port usage in the following examples is for an AFF A700 or a FAS9000 system.

Steps

1. Confirm that each node has disk automatic assignment enabled:

   ```
   storage disk option show
   ```

   Disk automatic assignment will assign pool 0 and pool 1 disks on a shelf-by-shelf basis.

   The Auto Assign column indicates whether disk automatic assignment is enabled.

   ```
   Node         BKg. FW. Upd.  Auto Copy   Auto Assign  Auto Assign Policy
   ----------   -------------  ----------  -----------  ------------------
   node_A_1             on          on           on              default
   node_A_2             on          on           on              default
   2 entries were displayed.
   ```

2. Verify you can create MetroCluster IP interfaces on the nodes:

```
metrocluster configuration-settings show-status
```

All nodes should be ready:

```
Cluster          Node          Configuration Settings Status
----------       -----------   --------------------------------
cluster_A
                 node_A_1      ready for interface create
                 node_A_2      ready for interface create
cluster_B
                 node_B_1      ready for interface create
                 node_B_2      ready for interface create
4 entries were displayed.
```

3. Create the interfaces on "node_A_1".

> (i)
> ◦ The port usage in the following examples is for an AFF A700 or a FAS9000 system (e5a and e5b). You must configure the interfaces on the correct ports for your platform model, as given above.
>
> ◦ Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also specify the -gateway parameter when creating MetroCluster IP interfaces. Refer to Considerations for layer 3 wide-area networks.
>
> ◦ On platform models that support VLANs for the MetroCluster IP interface, you can include the -vlan-id parameter if you don't want to use the default VLAN IDs.

a. Configure the interface on port "e5a" on "node_A_1":

```
metrocluster configuration-settings interface create -cluster-name cluster-name -home-node node-name -home-port e5a -address ip-address -netmask netmask
```

The following example shows the creation of the interface on port "e5a" on "node_A_1" with IP address "10.1.1.1":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_1 -home-port e5a -address
10.1.1.1 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

b. Configure the interface on port "e5b" on "node_A_1":

```
metrocluster configuration-settings interface create -cluster-name cluster-name -home-node node-name -home-port e5b -address ip-address -netmask netmask
```

The following example shows the creation of the interface on port "e5b" on "node_A_1" with IP address "10.1.2.1":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_1 -home-port e5b -address
10.1.2.1 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

(i) You can verify that these interfaces are present using the `metrocluster configuration-settings interface show` command.

4. Create the interfaces on "node_A_2".

(i)
   ◦ The port usage in the following examples is for an AFF A700 or a FAS9000 system ("e5a" and "e5b"). You must configure the interfaces on the correct ports for your platform model, as given above.
   ◦ Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also specify the `-gateway` parameter when creating MetroCluster IP interfaces. Refer to Considerations for layer 3 wide-area networks.
   ◦ On platform models that support VLANs for the MetroCluster IP interface, you can include the `-vlan-id` parameter if you don't want to use the default VLAN IDs.

a. Configure the interface on port "e5a" on "node_A_2":

```
metrocluster configuration-settings interface create -cluster-name cluster-
name -home-node node-name -home-port e5a -address ip-address -netmask
netmask
```

The following example shows the creation of the interface on port "e5a" on "node_A_2" with IP address "10.1.1.2":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_2 -home-port e5a -address
10.1.1.2 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

On platform models that support VLANs for the MetroCluster IP interface, you can include the `-vlan -id` paramter if you don't want to use the default VLAN IDs. The following example shows the command for an AFF A220 system with a VLAN ID of "120":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_2 -home-port e0a -address
10.1.1.2 -netmask 255.255.255.0 -vlan-id 120
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

b. Configure the interface on port "e5b" on "node_A_2":

```
metrocluster configuration-settings interface create -cluster-name cluster-
name -home-node node-name -home-port e5b -address ip-address -netmask
netmask
```

The following example shows the creation of the interface on port "e5b" on "node_A_2" with IP address "10.1.2.2":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_2 -home-port e5b -address
10.1.2.2 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

On platform models that support VLANs for the MetroCluster IP interface, you can include the `-vlan -id` paramter if you don't want to use the default VLAN IDs. The following example shows the command for an AFF A220 system with a VLAN ID of "220":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_2 -home-port e0b -address
10.1.2.2 -netmask 255.255.255.0 -vlan-id 220
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

5. Create the interfaces on "node_B_1".

(i)
- The port usage in the following examples is for an AFF A700 or a FAS9000 system ("e5a" and "e5b"). You must configure the interfaces on the correct ports for your platform model, as given above.

- Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also specify the `-gateway` parameter when creating MetroCluster IP interfaces. Refer to Considerations for layer 3 wide-area networks.

- On platform models that support VLANs for the MetroCluster IP interface, you can include the `-vlan-id` parameter if you don't want to use the default VLAN IDs.

a. Configure the interface on port "e5a" on "node_B_1":

```
metrocluster configuration-settings interface create -cluster-name cluster-
```

name -home-node *node-name* -home-port e5a -address *ip-address* -netmask
netmask

The following example shows the creation of the interface on port "e5a" on "node_B_1" with IP address
"10.1.1.3":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_B_1 -home-port e5a -address
10.1.1.3 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.cluster_A::>
```

b. Configure the interface on port "e5b" on "node_B_1":

metrocluster configuration-settings interface create -cluster-name *cluster-*
name -home-node *node-name* -home-port e5a -address *ip-address* -netmask
netmask

The following example shows the creation of the interface on port "e5b" on "node_B_1" with IP address
"10.1.2.3":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_B_1 -home-port e5b -address
10.1.2.3 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.cluster_A::>
```

6. Create the interfaces on "node_B_2".

> (i)
> - The port usage in the following examples is for an AFF A700 or a FAS9000 system (e5a
>   and e5b). You must configure the interfaces on the correct ports for your platform model,
>   as given above.
>
> - Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also
>   specify the -gateway parameter when creating MetroCluster IP interfaces. Refer to
>   Considerations for layer 3 wide-area networks.
>
> - On platform models that support VLANs for the MetroCluster IP interface, you can
>   include the -vlan-id parameter if you don't want to use the default VLAN IDs.

a. Configure the interface on port "e5a" on "node_B_2":

metrocluster configuration-settings interface create -cluster-name *cluster-*
name -home-node *node-name* -home-port e5a -address *ip-address* -netmask
netmask

The following example shows the creation of the interface on port "e5a" on "node_B_2" with IP address
"10.1.1.4":

```
cluster_B::>metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_2 -home-port e5a -address
10.1.1.4 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.cluster_A::>
```

b. Configure the interface on port "e5b" on "node_B_2":

```
metrocluster configuration-settings interface create -cluster-name cluster-
name -home-node node-name -home-port e5b -address ip-address -netmask
netmask
```

The following example shows the creation of the interface on port "e5b" on "node_B_2" with IP address "10.1.2.4":

```
cluster_B::> metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_2 -home-port e5b -address
10.1.2.4 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

7. Verify that the interfaces have been configured:

```
metrocluster configuration-settings interface show
```

The following example shows that the configuration state for each interface is completed.

```
cluster_A::> metrocluster configuration-settings interface show
DR                                                              Config
Group Cluster Node     Network Address Netmask        Gateway   State
----- ------- ------- --------------- --------------- ---------
----------
1     cluster_A  node_A_1
                Home Port: e5a
                      10.1.1.1     255.255.255.0   -         completed
                Home Port: e5b
                      10.1.2.1     255.255.255.0   -         completed
                node_A_2
                Home Port: e5a
                      10.1.1.2     255.255.255.0   -         completed
                Home Port: e5b
                      10.1.2.2     255.255.255.0   -         completed
      cluster_B  node_B_1
                Home Port: e5a
                      10.1.1.3     255.255.255.0   -         completed
                Home Port: e5b
                      10.1.2.3     255.255.255.0   -         completed
                node_B_2
                Home Port: e5a
                      10.1.1.4     255.255.255.0   -         completed
                Home Port: e5b
                      10.1.2.4     255.255.255.0   -         completed
8 entries were displayed.
cluster_A::>
```

8. Verify that the nodes are ready to connect the MetroCluster interfaces:

```
metrocluster configuration-settings show-status
```

The following example shows all nodes in the "ready for connection" state:

```
Cluster        Node         Configuration Settings Status
----------     -----------  --------------------------------
cluster_A
               node_A_1     ready for connection connect
               node_A_2     ready for connection connect
cluster_B
               node_B_1     ready for connection connect
               node_B_2     ready for connection connect
4 entries were displayed.
```

9. Establish the connections:

```
metrocluster configuration-settings connection connect
```

The IP addresses cannot be changed after you issue this command.

The following example shows "cluster_A" is successfully connected:

```
cluster_A::> metrocluster configuration-settings connection connect
[Job 53] Job succeeded: Connect is successful.
cluster_A::>
```

10. Verify that the connections have been established:

```
metrocluster configuration-settings show-status
```

The configuration settings status for all nodes should be completed:

```
Cluster          Node           Configuration Settings Status
----------       -----------    --------------------------------
cluster_A
                 node_A_1       completed
                 node_A_2       completed
cluster_B
                 node_B_1       completed
                 node_B_2       completed
4 entries were displayed.
```

11. Verify that the iSCSI connections have been established:

    a. Change to the advanced privilege level:

    ```
    set -privilege advanced
    ```

    You need to respond with "y" when you are prompted to continue into advanced mode and you see the advanced mode prompt (*>).

    b. Display the connections:

    ```
    storage iscsi-initiator show
    ```

    On systems running ONTAP 9.5, there are eight MetroCluster IP initiators on each cluster that should appear in the output.

    On systems running ONTAP 9.4 and earlier, there are four MetroCluster IP initiators on each cluster that should appear in the output.

    The following example shows the eight MetroCluster IP initiators on a cluster running ONTAP 9.5:

```
cluster_A::*> storage iscsi-initiator show
Node Type Label     Target Portal           Target Name
Admin/Op
---- ---- -------- -----------------
------------------------------ --------


cluster_A-01
    dr_auxiliary
              mccip-aux-a-initiator
                    10.227.16.113:65200     prod506.com.company:abab44
up/up
              mccip-aux-a-initiator2
                    10.227.16.113:65200     prod507.com.company:abab44
up/up
              mccip-aux-b-initiator
                    10.227.95.166:65200     prod506.com.company:abab44
up/up
              mccip-aux-b-initiator2
                    10.227.95.166:65200     prod507.com.company:abab44
up/up
    dr_partner
              mccip-pri-a-initiator
                    10.227.16.112:65200     prod506.com.company:cdcd88
up/up
              mccip-pri-a-initiator2
                    10.227.16.112:65200     prod507.com.company:cdcd88
up/up
              mccip-pri-b-initiator
                    10.227.95.165:65200     prod506.com.company:cdcd88
up/up
              mccip-pri-b-initiator2
                    10.227.95.165:65200     prod507.com.company:cdcd88
up/up
cluster_A-02
    dr_auxiliary
              mccip-aux-a-initiator
                    10.227.16.112:65200     prod506.com.company:cdcd88
up/up
              mccip-aux-a-initiator2
                    10.227.16.112:65200     prod507.com.company:cdcd88
up/up
              mccip-aux-b-initiator
                    10.227.95.165:65200     prod506.com.company:cdcd88
up/up
              mccip-aux-b-initiator2
                    10.227.95.165:65200     prod507.com.company:cdcd88
```

```
    up/up
         dr_partner
                  mccip-pri-a-initiator
                        10.227.16.113:65200      prod506.com.company:abab44
    up/up
                  mccip-pri-a-initiator2
                        10.227.16.113:65200      prod507.com.company:abab44
    up/up
                  mccip-pri-b-initiator
                        10.227.95.166:65200      prod506.com.company:abab44
    up/up
                  mccip-pri-b-initiator2
                        10.227.95.166:65200      prod507.com.company:abab44
    up/up
    16 entries were displayed.
```

c. Return to the admin privilege level:

    `set -privilege admin`

12. Verify that the nodes are ready for final implementation of the MetroCluster configuration:

`metrocluster node show`

```
cluster_A::> metrocluster node show
DR                                 Configuration  DR
Group Cluster Node                 State          Mirroring Mode
----- ------- ------------------ -------------- --------- ----
-     cluster_A
              node_A_1           ready to configure -      -
              node_A_2           ready to configure -      -
2 entries were displayed.
cluster_A::>
```

```
cluster_B::> metrocluster node show
DR                                 Configuration  DR
Group Cluster Node                 State          Mirroring Mode
----- ------- ------------------ -------------- --------- ----
-     cluster_B
              node_B_1           ready to configure -      -
              node_B_2           ready to configure -      -
2 entries were displayed.
cluster_B::>
```

Verifying or manually performing pool 1 drives assignment

Depending on the storage configuration, you must either verify pool 1 drive assignment or manually assign drives to pool 1 for each node in the MetroCluster IP configuration.

About this task

The procedure you use depends on the version of ONTAP you are using.

Configuration type	Procedure
The systems meet the requirements for automatic drive assignment or, if running ONTAP 9.3, were received from the factory.	Verifying disk assignment for pool 1 disks
The configuration includes either three shelves, or, if it contains more than four shelves, has an uneven multiple of four shelves (for example, seven shelves), and is running ONTAP 9.5.	Manually assigning drives for pool 1 (ONTAP 9.4 or later)
The configuration does not include four storage shelves per site and is running ONTAP 9.4	Manually assigning drives for pool 1 (ONTAP 9.4 or later)
The systems were not received from the factory and are running ONTAP 9.3Systems received from the factory are pre-configured with assigned drives.	Manually assigning disks for pool 1 (ONTAP 9.3)

Verifying disk assignment for pool 1 disks

You must verify that the remote disks are visible to the nodes and have been assigned correctly.

About this task

You must wait at least ten minutes for disk auto-assignment to complete after the MetroCluster IP interfaces and connections were created with the `metrocluster configuration-settings connection connect` command.

Command output will show disk names in the following format:

`node-name:0m.i1.0L1`

Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later

Step

1. Verify that pool 1 disks are auto-assigned:

   `disk show`

   The following output shows the output for an AFF A800 system with no external shelves.

   Drive autoassignment has assigned one quarter (8 drives) to "node_A_1" and one quarter to "node_A_2". The remaining drives will be remote (pool1) disks for "node_B_1" and "node_B_2".

```
cluster_B::> disk show -host-adapter 0m -owner node_B_2
                     Usable      Disk              Container  Container
Disk                 Size        Shelf Bay Type    Type       Name
Owner
----------------     ---------- ----- --- ------- ----------- ---------
--------
node_B_2:0m.i0.2L4   894.0GB     0     29  SSD-NVM shared      -
node_B_2
node_B_2:0m.i0.2L10  894.0GB     0     25  SSD-NVM shared      -
node_B_2
node_B_2:0m.i0.3L3   894.0GB     0     28  SSD-NVM shared      -
node_B_2
node_B_2:0m.i0.3L9   894.0GB     0     24  SSD-NVM shared      -
node_B_2
node_B_2:0m.i0.3L11  894.0GB     0     26  SSD-NVM shared      -
node_B_2
node_B_2:0m.i0.3L12  894.0GB     0     27  SSD-NVM shared      -
node_B_2
node_B_2:0m.i0.3L15  894.0GB     0     30  SSD-NVM shared      -
node_B_2
node_B_2:0m.i0.3L16  894.0GB     0     31  SSD-NVM shared      -
node_B_2
8 entries were displayed.

cluster_B::> disk show -host-adapter 0m -owner node_B_1
                     Usable      Disk              Container  Container
Disk                 Size        Shelf Bay Type    Type       Name
Owner
----------------     ---------- ----- --- ------- ----------- ---------
--------
node_B_1:0m.i2.3L19  1.75TB      0     42  SSD-NVM shared      -
node_B_1
node_B_1:0m.i2.3L20  1.75TB      0     43  SSD-NVM spare       Pool1
node_B_1
node_B_1:0m.i2.3L23  1.75TB      0     40  SSD-NVM shared      -
node_B_1
node_B_1:0m.i2.3L24  1.75TB      0     41  SSD-NVM spare       Pool1
node_B_1
node_B_1:0m.i2.3L29  1.75TB      0     36  SSD-NVM shared      -
node_B_1
node_B_1:0m.i2.3L30  1.75TB      0     37  SSD-NVM shared      -
node_B_1
node_B_1:0m.i2.3L31  1.75TB      0     38  SSD-NVM shared      -
node_B_1
node_B_1:0m.i2.3L32  1.75TB      0     39  SSD-NVM shared      -
node_B_1
```

```
8 entries were displayed.

cluster_B::> disk show
                     Usable      Disk                Container   Container
Disk                 Size        Shelf Bay Type      Type        Name
Owner
----------------     ----------  ----- --- -------   ----------- ---------
--------
node_B_1:0m.i1.0L6   1.75TB      0     1   SSD-NVM shared        -
node_A_2
node_B_1:0m.i1.0L8   1.75TB      0     3   SSD-NVM shared        -
node_A_2
node_B_1:0m.i1.0L17  1.75TB      0     18  SSD-NVM shared        -
node_A_1
node_B_1:0m.i1.0L22  1.75TB      0     17 SSD-NVM shared - node_A_1
node_B_1:0m.i1.0L25  1.75TB      0     12 SSD-NVM shared - node_A_1
node_B_1:0m.i1.2L2   1.75TB      0     5 SSD-NVM shared - node_A_2
node_B_1:0m.i1.2L7   1.75TB      0     2 SSD-NVM shared - node_A_2
node_B_1:0m.i1.2L14  1.75TB      0     7 SSD-NVM shared - node_A_2
node_B_1:0m.i1.2L21  1.75TB      0     16 SSD-NVM shared - node_A_1
node_B_1:0m.i1.2L27  1.75TB      0     14 SSD-NVM shared - node_A_1
node_B_1:0m.i1.2L28  1.75TB      0     15 SSD-NVM shared - node_A_1
node_B_1:0m.i2.1L1   1.75TB      0     4 SSD-NVM shared - node_A_2
node_B_1:0m.i2.1L5   1.75TB      0     0 SSD-NVM shared - node_A_2
node_B_1:0m.i2.1L13  1.75TB      0     6 SSD-NVM shared - node_A_2
node_B_1:0m.i2.1L18  1.75TB      0     19 SSD-NVM shared - node_A_1
node_B_1:0m.i2.1L26  1.75TB      0     13 SSD-NVM shared - node_A_1
node_B_1:0m.i2.3L19  1.75TB      0 42 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L20  1.75TB      0 43 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L23  1.75TB      0 40 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L24  1.75TB      0 41 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L29  1.75TB      0 36 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L30  1.75TB      0 37 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L31  1.75TB      0 38 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L32  1.75TB      0 39 SSD-NVM shared - node_B_1
node_B_1:0n.12       1.75TB      0 12 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.13       1.75TB      0 13 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.14       1.75TB      0 14 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.15       1.75TB 0 15 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.16       1.75TB 0 16 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.17       1.75TB 0 17 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.18       1.75TB 0 18 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.19       1.75TB 0 19 SSD-NVM shared - node_B_1
node_B_1:0n.24       894.0GB 0 24 SSD-NVM shared - node_A_2
node_B_1:0n.25       894.0GB 0 25 SSD-NVM shared - node_A_2
node_B_1:0n.26       894.0GB 0 26 SSD-NVM shared - node_A_2
```

```
node_B_1:0n.27         894.0GB 0 27 SSD-NVM shared - node_A_2
node_B_1:0n.28         894.0GB 0 28 SSD-NVM shared - node_A_2
node_B_1:0n.29         894.0GB 0 29 SSD-NVM shared - node_A_2
node_B_1:0n.30         894.0GB 0 30 SSD-NVM shared - node_A_2
node_B_1:0n.31         894.0GB 0 31 SSD-NVM shared - node_A_2
node_B_1:0n.36         1.75TB 0 36 SSD-NVM shared - node_A_1
node_B_1:0n.37         1.75TB 0 37 SSD-NVM shared - node_A_1
node_B_1:0n.38         1.75TB 0 38 SSD-NVM shared - node_A_1
node_B_1:0n.39         1.75TB 0 39 SSD-NVM shared - node_A_1
node_B_1:0n.40         1.75TB 0 40 SSD-NVM shared - node_A_1
node_B_1:0n.41         1.75TB 0 41 SSD-NVM shared - node_A_1
node_B_1:0n.42         1.75TB 0 42 SSD-NVM shared - node_A_1
node_B_1:0n.43         1.75TB 0 43 SSD-NVM shared - node_A_1
node_B_2:0m.i0.2L4  894.0GB 0 29 SSD-NVM shared - node_B_2
node_B_2:0m.i0.2L10 894.0GB 0 25 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L3  894.0GB 0 28 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L9  894.0GB 0 24 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L11 894.0GB 0 26 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L12 894.0GB 0 27 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L15 894.0GB 0 30 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L16 894.0GB 0 31 SSD-NVM shared - node_B_2
node_B_2:0n.0       1.75TB 0 0 SSD-NVM shared aggr0_rha12_b1_cm_02_0
node_B_2
node_B_2:0n.1 1.75TB 0 1 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.2 1.75TB 0 2 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.3 1.75TB 0 3 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.4 1.75TB 0 4 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.5 1.75TB 0 5 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.6 1.75TB 0 6 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.7 1.75TB 0 7 SSD-NVM shared - node_B_2
64 entries were displayed.


cluster_B::>



cluster_A::> disk show
Usable Disk Container Container
Disk Size Shelf Bay Type Type Name Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------
node_A_1:0m.i1.0L2 1.75TB 0 5 SSD-NVM shared - node_B_2
node_A_1:0m.i1.0L8 1.75TB 0 3 SSD-NVM shared - node_B_2
node_A_1:0m.i1.0L18 1.75TB 0 19 SSD-NVM shared - node_B_1
node_A_1:0m.i1.0L25 1.75TB 0 12 SSD-NVM shared - node_B_1
node_A_1:0m.i1.0L27 1.75TB 0 14 SSD-NVM shared - node_B_1
node_A_1:0m.i1.2L1 1.75TB 0 4 SSD-NVM shared - node_B_2
```

```
node_A_1:0m.i1.2L6 1.75TB 0 1 SSD-NVM shared - node_B_2
node_A_1:0m.i1.2L7 1.75TB 0 2 SSD-NVM shared - node_B_2
node_A_1:0m.i1.2L14 1.75TB 0 7 SSD-NVM shared - node_B_2
node_A_1:0m.i1.2L17 1.75TB 0 18 SSD-NVM shared - node_B_1
node_A_1:0m.i1.2L22 1.75TB 0 17 SSD-NVM shared - node_B_1
node_A_1:0m.i2.1L5 1.75TB 0 0 SSD-NVM shared - node_B_2
node_A_1:0m.i2.1L13 1.75TB 0 6 SSD-NVM shared - node_B_2
node_A_1:0m.i2.1L21 1.75TB 0 16 SSD-NVM shared - node_B_1
node_A_1:0m.i2.1L26 1.75TB 0 13 SSD-NVM shared - node_B_1
node_A_1:0m.i2.1L28 1.75TB 0 15 SSD-NVM shared - node_B_1
node_A_1:0m.i2.3L19 1.75TB 0 42 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L20 1.75TB 0 43 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L23 1.75TB 0 40 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L24 1.75TB 0 41 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L29 1.75TB 0 36 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L30 1.75TB 0 37 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L31 1.75TB 0 38 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L32 1.75TB 0 39 SSD-NVM shared - node_A_1
node_A_1:0n.12 1.75TB 0 12 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.13 1.75TB 0 13 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.14 1.75TB 0 14 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.15 1.75TB 0 15 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.16 1.75TB 0 16 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.17 1.75TB 0 17 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.18 1.75TB 0 18 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.19 1.75TB 0 19 SSD-NVM shared - node_A_1
node_A_1:0n.24 894.0GB 0 24 SSD-NVM shared - node_B_2
node_A_1:0n.25 894.0GB 0 25 SSD-NVM shared - node_B_2
node_A_1:0n.26 894.0GB 0 26 SSD-NVM shared - node_B_2
node_A_1:0n.27 894.0GB 0 27 SSD-NVM shared - node_B_2
node_A_1:0n.28 894.0GB 0 28 SSD-NVM shared - node_B_2
node_A_1:0n.29 894.0GB 0 29 SSD-NVM shared - node_B_2
node_A_1:0n.30 894.0GB 0 30 SSD-NVM shared - node_B_2
node_A_1:0n.31 894.0GB 0 31 SSD-NVM shared - node_B_2
node_A_1:0n.36 1.75TB 0 36 SSD-NVM shared - node_B_1
node_A_1:0n.37 1.75TB 0 37 SSD-NVM shared - node_B_1
node_A_1:0n.38 1.75TB 0 38 SSD-NVM shared - node_B_1
node_A_1:0n.39 1.75TB 0 39 SSD-NVM shared - node_B_1
node_A_1:0n.40 1.75TB 0 40 SSD-NVM shared - node_B_1
node_A_1:0n.41 1.75TB 0 41 SSD-NVM shared - node_B_1
node_A_1:0n.42 1.75TB 0 42 SSD-NVM shared - node_B_1
node_A_1:0n.43 1.75TB 0 43 SSD-NVM shared - node_B_1
node_A_2:0m.i2.3L3 894.0GB 0 28 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L4 894.0GB 0 29 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L9 894.0GB 0 24 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L10 894.0GB 0 25 SSD-NVM shared - node_A_2
```

```
node_A_2:0m.i2.3L11 894.0GB 0 26 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L12 894.0GB 0 27 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L15 894.0GB 0 30 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L16 894.0GB 0 31 SSD-NVM shared - node_A_2
node_A_2:0n.0 1.75TB 0 0 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.1 1.75TB 0 1 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.2 1.75TB 0 2 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.3 1.75TB 0 3 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.4 1.75TB 0 4 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.5 1.75TB 0 5 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.6 1.75TB 0 6 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.7 1.75TB 0 7 SSD-NVM shared - node_A_2
64 entries were displayed.

cluster_A::>
```

Manually assigning drives for pool 1 (ONTAP 9.4 or later)

If the system was not preconfigured at the factory and does not meet the requirements for automatic drive assignment, you must manually assign the remote pool 1 drives.

About this task

This procedure applies to configurations running ONTAP 9.4 or later.

Details for determining whether your system requires manual disk assignment are included in Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later.

When the configuration includes only two external shelves per site, pool 1 drives for each site should be shared from the same shelf as shown in the following examples:

- node_A_1 is assigned drives in bays 0-11 on site_B-shelf_2 (remote)
- node_A_2 is assigned drives in bays 12-23 on site_B-shelf_2 (remote)

Steps

1. From each node in the MetroCluster IP configuration, assign remote drives to pool 1.

   a. Display the list of unassigned drives:

   ```
   disk show -host-adapter 0m -container-type unassigned
   ```

```
cluster_A::> disk show -host-adapter 0m -container-type unassigned
                     Usable              Disk    Container   Container
Disk                 Size Shelf Bay Type Type        Name
Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------
6.23.0                  -    23    0 SSD     unassigned  -           -
6.23.1                  -    23    1 SSD     unassigned  -           -
.
.
.
node_A_2:0m.i1.2L51     -    21   14 SSD     unassigned  -           -
node_A_2:0m.i1.2L64     -    21   10 SSD     unassigned  -           -
.
.
.
48 entries were displayed.

cluster_A::>
```

b. Assign ownership of remote drives (0m) to pool 1 of the first node (for example, "node_A_1"):

```
disk assign -disk disk-id -pool 1 -owner owner-node-name
```

The `disk-id` must identify a drive on a remote shelf of `owner-node-name`.

c. Confirm that the drives were assigned to pool 1:

```
disk show -host-adapter 0m -container-type unassigned
```

(i) The iSCSI connection used to access the remote drives appears as device 0m.

The following output shows that the drives on shelf "23" were assigned because they no longer appear in the list of unassigned drives:

```
cluster_A::> disk show -host-adapter 0m -container-type unassigned
                        Usable              Disk    Container   Container
Disk                     Size Shelf Bay Type    Type        Name
Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------
node_A_2:0m.i1.2L51        -    21  14 SSD     unassigned  -          -
node_A_2:0m.i1.2L64        -    21  10 SSD     unassigned  -          -
.
.
.
node_A_2:0m.i2.1L90        -    21  19 SSD     unassigned  -          -
24 entries were displayed.

cluster_A::>
```

    d. Repeat these steps to assign pool 1 drives to the second node on site A (for example, "node_A_2").

    e. Repeat these steps on site B.

Manually assigning disks for pool 1 (ONTAP 9.3)

If you have at least two disk shelves for each node, you use ONTAP's auto-assignment functionality to automatically assign the remote (pool 1) disks.

About this task

You must first assign a disk on the shelf to pool 1. ONTAP then automatically assigns the rest of the disks on the shelf to the same pool.

This procedure applies to configurations running ONTAP 9.3.

This procedure can be used only if you have at least two disk shelves for each node, which allows shelf-level auto-assignment of disks.

If you cannot use shelf-level auto-assignment, you must manually assign your remote disks so that each node has a remote pool of disks (pool 1).

The ONTAP automatic disk assignment feature assigns the disks on a shelf-by-shelf basis. For example:

- All the disks on site_B-shelf_2 are autoassigned to pool1 of node_A_1
- All the disks on site_B-shelf_4 are autoassigned to pool1 of node_A_2
- All the disks on site_A-shelf_2 are autoassigned to pool1 of node_B_1
- All the disks on site_A-shelf_4 are autoassigned to pool1 of node_B_2

You must "seed" the auto-assignment by specifying a single disk on each shelf.

Steps

1. From each node in the MetroCluster IP configuration, assign a remote disk to pool 1.

    a. Display the list of unassigned disks:

```
disk show -host-adapter 0m -container-type unassigned
```

```
cluster_A::> disk show -host-adapter 0m -container-type unassigned
                        Usable           Disk    Container   Container
Disk                    Size Shelf Bay Type    Type        Name
Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------
6.23.0                     -    23   0 SSD     unassigned  -          -
6.23.1                     -    23   1 SSD     unassigned  -          -
.
.
.
node_A_2:0m.i1.2L51        -    21  14 SSD     unassigned  -          -
node_A_2:0m.i1.2L64        -    21  10 SSD     unassigned  -          -
.
.
.
48 entries were displayed.

cluster_A::>
```

b. Select a remote disk (0m) and assign ownership of the disk to pool 1 of the first node (for example, "node_A_1"):

```
disk assign -disk disk-id -pool 1 -owner owner-node-name
```

The `disk-id` must identify a disk on a remote shelf of `owner-node-name`.

The ONTAP disk auto-assignment feature assigns all disks on the remote shelf that contains the specified disk.

c. After waiting at least 60 seconds for disk auto-assignment to take place, verify that the remote disks on the shelf were auto-assigned to pool 1:

```
disk show -host-adapter 0m -container-type unassigned
```

ⓘ The iSCSI connection used to access the remote disks appears as device 0m.

The following output shows that the disks on shelf "23" have now been assigned and no longer appear:

```
cluster_A::> disk show -host-adapter 0m -container-type unassigned
                          Usable                Disk      Container   Container
Disk                      Size Shelf Bay Type   Type        Name
Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------
node_A_2:0m.i1.2L51       -    21  14 SSD      unassigned  -           -
node_A_2:0m.i1.2L64       -    21  10 SSD      unassigned  -           -
node_A_2:0m.i1.2L72       -    21  23 SSD      unassigned  -           -
node_A_2:0m.i1.2L74       -    21   1 SSD      unassigned  -           -
node_A_2:0m.i1.2L83       -    21  22 SSD      unassigned  -           -
node_A_2:0m.i1.2L90       -    21   7 SSD      unassigned  -           -
node_A_2:0m.i1.3L52       -    21   6 SSD      unassigned  -           -
node_A_2:0m.i1.3L59       -    21  13 SSD      unassigned  -           -
node_A_2:0m.i1.3L66       -    21  17 SSD      unassigned  -           -
node_A_2:0m.i1.3L73       -    21  12 SSD      unassigned  -           -
node_A_2:0m.i1.3L80       -    21   5 SSD      unassigned  -           -
node_A_2:0m.i1.3L81       -    21   2 SSD      unassigned  -           -
node_A_2:0m.i1.3L82       -    21  16 SSD      unassigned  -           -
node_A_2:0m.i1.3L91       -    21   3 SSD      unassigned  -           -
node_A_2:0m.i2.0L49       -    21  15 SSD      unassigned  -           -
node_A_2:0m.i2.0L50       -    21   4 SSD      unassigned  -           -
node_A_2:0m.i2.1L57       -    21  18 SSD      unassigned  -           -
node_A_2:0m.i2.1L58       -    21  11 SSD      unassigned  -           -
node_A_2:0m.i2.1L59       -    21  21 SSD      unassigned  -           -
node_A_2:0m.i2.1L65       -    21  20 SSD      unassigned  -           -
node_A_2:0m.i2.1L72       -    21   9 SSD      unassigned  -           -
node_A_2:0m.i2.1L80       -    21   0 SSD      unassigned  -           -
node_A_2:0m.i2.1L88       -    21   8 SSD      unassigned  -           -
node_A_2:0m.i2.1L90       -    21  19 SSD      unassigned  -           -
24 entries were displayed.

cluster_A::>
```

d. Repeat these steps to assign pool 1 disks to the second node on site A (for example, "node_A_2").

e. Repeat these steps on site B.

Enabling automatic drive assignment in ONTAP 9.4

In ONTAP 9.4, if you disabled automatic drive assignment as directed previously in this procedure, you must reenable it on all nodes.

Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later

Step

1. Enable automatic drive assignment:

```
storage disk option modify -node node_name -autoassign on
```

You must issue this command on all nodes in the MetroCluster IP configuration.

Mirroring the root aggregates

You must mirror the root aggregates to provide data protection.

About this task

By default, the root aggregate is created as RAID-DP type aggregate. You can change the root aggregate from RAID-DP to RAID4 type aggregate. The following command modifies the root aggregate for RAID4 type aggregate:

```
storage aggregate modify –aggregate aggr_name -raidtype raid4
```

> ⓘ   On non-ADP systems, the RAID type of the aggregate can be modified from the default RAID-DP to RAID4 before or after the aggregate is mirrored.

Steps

1. Mirror the root aggregate:

   ```
   storage aggregate mirror aggr_name
   ```

   The following command mirrors the root aggregate for controller_A_1:

   ```
   controller_A_1::> storage aggregate mirror aggr0_controller_A_1
   ```

   This mirrors the aggregate, so it consists of a local plex and a remote plex located at the remote MetroCluster site.

2. Repeat the previous step for each node in the MetroCluster configuration.

Related information

Logical storage management

Creating a mirrored data aggregate on each node

You must create a mirrored data aggregate on each node in the DR group.

About this task

- You should know what drives will be used in the new aggregate.
- If you have multiple drive types in your system (heterogeneous storage), you should understand how you can ensure that the correct drive type is selected.
- Drives are owned by a specific node; when you create an aggregate, all drives in that aggregate must be owned by the same node, which becomes the home node for that aggregate.

  In systems using ADP, aggregates are created using partitions in which each drive is partitioned in to P1, P2 and P3 partitions.

- Aggregate names should conform to the naming scheme you determined when you planned your

MetroCluster configuration.

Steps

1. Display a list of available spares:

   ```
   storage disk show -spare -owner node_name
   ```

2. Create the aggregate:

   ```
   storage aggregate create -mirror true
   ```

   If you are logged in to the cluster on the cluster management interface, you can create an aggregate on any node in the cluster. To ensure that the aggregate is created on a specific node, use the `-node `parameter or specify drives that are owned by that node.

   You can specify the following options:

   ◦ Aggregate's home node (that is, the node that owns the aggregate in normal operation).
   ◦ List of specific drives that are to be added to the aggregate.
   ◦ Number of drives to include.

   > (i) In the minimum supported configuration, in which a limited number of drives are available, you must use the force-small-aggregate option to allow the creation of a three disk RAID-DP aggregate.

   ◦ Checksum style to use for the aggregate.
   ◦ Type of drives to use.
   ◦ Size of drives to use.
   ◦ Drive speed to use.
   ◦ RAID type for RAID groups on the aggregate.
   ◦ Maximum number of drives that can be included in a RAID group.
   ◦ Whether drives with different RPM are allowed.

   For more information about these options, see the storage aggregate create man page.

   The following command creates a mirrored aggregate with 10 disks:

   ```
   cluster_A::> storage aggregate create aggr1_node_A_1 -diskcount 10
   -node node_A_1 -mirror true
   [Job 15] Job is queued: Create aggr1_node_A_1.
   [Job 15] The job is starting.
   [Job 15] Job succeeded: DONE
   ```

3. Verify the RAID group and drives of your new aggregate:

```
storage aggregate show-status -aggregate aggregate-name
```

Implementing the MetroCluster configuration

You must run the `metrocluster configure` command to start data protection in a MetroCluster configuration.

About this task

- There should be at least two non-root mirrored data aggregates on each cluster.

  You can verify this with the storage aggregate show command.

  > (i) If you want to use a single mirrored data aggregate, then see Step 1 for instructions.

- The ha-config state of the controllers and chassis must be "mccip".

You issue the `metrocluster configure` command once, on any of the nodes, to enable the MetroCluster configuration. You do not need to issue the command on each of the sites or nodes, and it does not matter which node or site you choose to issue the command on.

The `metrocluster configure` command automatically pairs the two nodes with the lowest system IDs in each of the two clusters as disaster recovery (DR) partners. In a four-node MetroCluster configuration, there are two DR partner pairs. The second DR pair is created from the two nodes with higher system IDs.

Steps

1. Configure the MetroCluster in the following format:

If your MetroCluster configuration has…	Then do this…
Multiple data aggregates	From any node's prompt, configure MetroCluster: `metrocluster configure node-name`

A single mirrored data aggregate	a. From any node's prompt, change to the advanced privilege level:
	
	`set -privilege advanced`
	
	You need to respond with "y" when you are prompted to continue into advanced mode and you see the advanced mode prompt (`*>`).
	
	b. Configure the MetroCluster with the `-allow -with-one-aggregate` true parameter:
	
	`metrocluster configure -allow-with -one-aggregate true node-name`
	
	c. Return to the admin privilege level:
	
	`set -privilege admin`

> **ⓘ** The best practice is to have multiple data aggregates. If the first DR group has only one aggregate and you want to add a DR group with one aggregate, you must move the metadata volume off the single data aggregate. For more information on this procedure, see Moving a metadata volume in MetroCluster configurations.

The following command enables the MetroCluster configuration on all of the nodes in the DR group that contains "controller_A_1":

```
cluster_A::*> metrocluster configure -node-name controller_A_1

[Job 121] Job succeeded: Configure is successful.
```

2. Verify the networking status on site A:

```
network port show
```

The following example shows the network port usage on a four-node MetroCluster configuration:

```
cluster_A::> network port show
                                                      Speed (Mbps)
Node    Port       IPspace    Broadcast Domain Link   MTU    Admin/Oper
------  ---------  ---------  ---------------- -----  ------- ------------
controller_A_1
        e0a        Cluster    Cluster          up     9000   auto/1000
        e0b        Cluster    Cluster          up     9000   auto/1000
        e0c        Default    Default          up     1500   auto/1000
        e0d        Default    Default          up     1500   auto/1000
        e0e        Default    Default          up     1500   auto/1000
        e0f        Default    Default          up     1500   auto/1000
        e0g        Default    Default          up     1500   auto/1000
controller_A_2
        e0a        Cluster    Cluster          up     9000   auto/1000
        e0b        Cluster    Cluster          up     9000   auto/1000
        e0c        Default    Default          up     1500   auto/1000
        e0d        Default    Default          up     1500   auto/1000
        e0e        Default    Default          up     1500   auto/1000
        e0f        Default    Default          up     1500   auto/1000
        e0g        Default    Default          up     1500   auto/1000
14 entries were displayed.
```

3. Verify the MetroCluster configuration from both sites in the MetroCluster configuration.

   a. Verify the configuration from site A:

```
metrocluster show
```

```
cluster_A::> metrocluster show

Configuration: IP fabric

Cluster                       Entry Name          State
------------------------       -------------------  -----------
 Local: cluster_A              Configuration state configured
                               Mode                normal
Remote: cluster_B              Configuration state configured
                               Mode                normal
```

b. Verify the configuration from site B:

`metrocluster show`

```
cluster_B::> metrocluster show

Configuration: IP fabric

Cluster                       Entry Name          State
------------------------       -------------------  -----------
 Local: cluster_B              Configuration state configured
                               Mode                normal
Remote: cluster_A              Configuration state configured
                               Mode                normal
```

4. To avoid possible issues with nonvolatile memory mirroring, reboot each of the four nodes:

`node reboot -node node-name -inhibit-takeover true`

5. Issue the `metrocluster show` command on both clusters to again verify the configuration.

Configuring the second DR group in an eight-node configuration

Repeat the previous tasks to configure the nodes in the second DR group.

Creating unmirrored data aggregates

You can optionally create unmirrored data aggregates for data that does not require the redundant mirroring provided by MetroCluster configurations.

About this task

- You should know what drives or array LUNs will be used in the new aggregate.

- If you have multiple drive types in your system (heterogeneous storage), you should understand how you can verify that the correct drive type is selected.

 In MetroCluster IP configurations, remote unmirrored aggregates are not accessible after a switchover

 The unmirrored aggregates must be local to the node owning them.

- Drives and array LUNs are owned by a specific node; when you create an aggregate, all drives in that aggregate must be owned by the same node, which becomes the home node for that aggregate.
- Aggregate names should conform to the naming scheme you determined when you planned your MetroCluster configuration.
- *Disks and aggregates management* contains more information about mirroring aggregates.

Steps

1. Enable unmirrored aggregate deployment:

   ```
   metrocluster modify -enable-unmirrored-aggr-deployment true
   ```

2. Verify that disk auto-assignment is disabled:

   ```
   disk option show
   ```

3. Install and cable the disk shelves that will contain the unmirrored aggregates.

   You can use the procedures in the *Installation and Setup* documentation for your platform and disk shelves.

   AFF and FAS Documentation Center

4. Manually assign all disks on the new shelf to the appropriate node:

   ```
   disk assign -disk disk-id -owner owner-node-name
   ```

5. Create the aggregate:

   ```
   storage aggregate create
   ```

   If you are logged in to the cluster on the cluster management interface, you can create an aggregate on any node in the cluster. To verify that the aggregate is created on a specific node, you should use the -node parameter or specify drives that are owned by that node.

   You must also ensure that you are only including drives on the unmirrored shelf to the aggregate.

   You can specify the following options:

   - Aggregate's home node (that is, the node that owns the aggregate in normal operation).
   - List of specific drives or array LUNs that are to be added to the aggregate.
   - Number of drives to include.
   - Checksum style to use for the aggregate.
   - Type of drives to use.
   - Size of drives to use.
   - Drive speed to use.

- RAID type for RAID groups on the aggregate.
- Maximum number of drives or array LUNs that can be included in a RAID group.
- Whether drives with different RPM are allowed.

  For more information about these options, see the storage aggregate create man page.

  The following command creates a unmirrored aggregate with 10 disks:

```
controller_A_1::> storage aggregate create aggr1_controller_A_1
-diskcount 10 -node controller_A_1
[Job 15] Job is queued: Create aggr1_controller_A_1.
[Job 15] The job is starting.
[Job 15] Job succeeded: DONE
```

6. Verify the RAID group and drives of your new aggregate:

   `storage aggregate show-status -aggregate aggregate-name`

7. Disable unmirrored aggregate deployment:

   `metrocluster modify -enable-unmirrored-aggr-deployment false`

8. Verify that disk autoassignment is enabled:

   `disk option show`

Related information

Disk and aggregate management

Checking the MetroCluster configuration

You can check that the components and relationships in the MetroCluster configuration are working correctly. You should do a check after initial configuration and after making any changes to the MetroCluster configuration. You should also do a check before a negotiated (planned) switchover or a switchback operation.

About this task

If the `metrocluster check run` command is issued twice within a short time on either or both clusters, a conflict can occur and the command might not collect all data. Subsequent `metrocluster check show` commands do not show the expected output.

Steps

1. Check the configuration:

   `metrocluster check run`

   The command runs as a background job and might not be completed immediately.

```
cluster_A::> metrocluster check run
The operation has been started and is running in the background. Wait
for
it to complete and run "metrocluster check show" to view the results. To
check the status of the running metrocluster check operation, use the
command,
"metrocluster operation history show -job-id 2245"
```

```
cluster_A::> metrocluster check show
Last Checked On: 9/13/2018 20:41:37

Component            Result
------------------- ---------
nodes                ok
lifs                 ok
config-replication   ok
aggregates           ok
clusters             ok
connections          ok
6 entries were displayed.
```

2. Display more detailed results from the most recent `metrocluster check run` command:

   `metrocluster check aggregate show`

   `metrocluster check cluster show`

   `metrocluster check config-replication show`

   `metrocluster check lif show`

   `metrocluster check node show`

   The `metrocluster check show` commands show the results of the most recent `metrocluster check run` command. You should always run the `metrocluster check run` command prior to using the `metrocluster check show` commands so that the information displayed is current.

   The following example shows the `metrocluster check aggregate show` command output for a healthy four-node MetroCluster configuration:

```
cluster_A::> metrocluster check aggregate show

Last Checked On: 8/5/2014 00:42:58

Node                    Aggregate                    Check
```

```
Result
--------------      ------------------      --------------------
---------
controller_A_1      controller_A_1_aggr0
                                            mirroring-status
ok
                                            disk-pool-allocation
ok
                                            ownership-state
ok
                    controller_A_1_aggr1
                                            mirroring-status
ok
                                            disk-pool-allocation
ok
                                            ownership-state
ok
                    controller_A_1_aggr2
                                            mirroring-status
ok
                                            disk-pool-allocation
ok
                                            ownership-state
ok


controller_A_2      controller_A_2_aggr0
                                            mirroring-status
ok
                                            disk-pool-allocation
ok
                                            ownership-state
ok
                    controller_A_2_aggr1
                                            mirroring-status
ok
                                            disk-pool-allocation
ok
                                            ownership-state
ok
                    controller_A_2_aggr2
                                            mirroring-status
ok
                                            disk-pool-allocation
ok
                                            ownership-state
```

```
ok


18 entries were displayed.
```

The following example shows the `metrocluster check cluster show` command output for a healthy four-node MetroCluster configuration. It indicates that the clusters are ready to perform a negotiated switchover if necessary.

```
Last Checked On: 9/13/2017 20:47:04

Cluster                  Check                           Result
-------------------- ---------------------------- ---------
mccint-fas9000-0102
                         negotiated-switchover-ready     not-applicable
                         switchback-ready                not-applicable
                         job-schedules                   ok
                         licenses                        ok
                         periodic-check-enabled          ok
mccint-fas9000-0304
                         negotiated-switchover-ready     not-applicable
                         switchback-ready                not-applicable
                         job-schedules                   ok
                         licenses                        ok
                         periodic-check-enabled          ok
10 entries were displayed.
```

Related information

Disk and aggregate management

Network and LIF management

Completing ONTAP configuration

After configuring, enabling, and checking the MetroCluster configuration, you can proceed to complete the cluster configuration by adding additional SVMs, network interfaces and other ONTAP functionality as needed.

Verifying switchover, healing, and switchback

You should verify the switchover, healing, and switchback operations of the MetroCluster configuration.

1. Use the procedures for negotiated switchover, healing, and switchback in MetroCluster management and disaster recovery

Configuring the MetroCluster Tiebreaker or ONTAP Mediator software

You can download and install on a third site either the MetroCluster Tiebreaker software, or beginning with ONTAP 9.7, the ONTAP Mediator.

Before you begin

You must have a Linux host available that has network connectivity to both clusters in the MetroCluster configuration. The specific requirements are in the MetroCluster Tiebreaker or ONTAP Mediator documentation.

If you are connecting to an existing Tiebreaker or ONTAP Mediator instance, you need the username, password, and IP address of the Tiebreaker or Mediator service.

If you must install a new instance of the ONTAP Mediator, follow the directions to install and configure the software.

[Configuring the ONTAP Mediator service for unplanned automatic switchover](#)

If you must install a new instance of the Tiebreaker software, follow the directions to install and configure the software.

[MetroCluster Tiebreaker Software installation and configuration](#)

About this task

You cannot use both the MetroCluster Tiebreaker software and the ONTAP Mediator with the same MetroCluster configuration.

[Considerations for using ONTAP Mediator or MetroCluster Tiebreaker](#)

Steps

1. Configure the ONTAP Mediator service or the Tiebreaker software:

   ◦ If you are using an existing instance of the ONTAP Mediator, add the ONTAP Mediator service to ONTAP:

     ```
     metrocluster configuration-settings mediator add -mediator-address ip-
     address-of-mediator-host
     ```

   ◦ If you are using the Tiebreaker software, refer to the Tiebreaker documentation.

     [MetroCluster Tiebreaker Software installation and configuration](#)

Protecting configuration backup files

You can provide additional protection for the cluster configuration backup files by specifying a remote URL (either HTTP or FTP) where the configuration backup files will be uploaded in addition to the default locations in the local cluster.

Steps

1. Set the URL of the remote destination for the configuration backup files:

   ```
   system configuration backup settings modify URL-of-destination
   ```

   Cluster Management with the CLI contains additional information under the section *Managing configuration backups*.

Related information

[System administration](#)

Restoring system defaults on a controller module

Reset and restore defaults on the controller modules.

1. At the LOADER prompt, return environmental variables to their default setting: `set-defaults`

2. Boot the node to the boot menu: `boot_ontap menu`

   After you run this command, wait until the boot menu is shown.

3. Clear the node configuration:

   ◦ If you are using systems configured for ADP, select option `9a` from the boot menu, and respond `yes` when prompted.

     ⓘ | This process is disruptive.

   The following screen shows the boot menu prompt:

   ```
   Please choose one of the following:

       (1) Normal Boot.
       (2) Boot without /etc/rc.
       (3) Change password.
       (4) Clean configuration and initialize all disks.
       (5) Maintenance mode boot.
       (6) Update flash from backup config.
       (7) Install new software first.
       (8) Reboot node.
       (9) Configure Advanced Drive Partitioning.
       Selection (1-9)?  9a
   ########## WARNING ##########

       This is a disruptive operation and will result in the
       loss of all filesystem data. Before proceeding further,
       make sure that:
       1) This option (9a) has been executed or will be executed
       on the HA partner node, prior to reinitializing either
       system in the HA-pair.
       2) The HA partner node is currently in a halted state or
       at the LOADER prompt.


       Do you still want to continue (yes/no)? yes
   ```

   ◦ If your system is not configured for ADP, type `wipeconfig` at the boot menu prompt, and then press Enter.

The following screen shows the boot menu prompt:

```
Please choose one of the following:

    (1) Normal Boot.
    (2) Boot without /etc/rc.
    (3) Change password.
    (4) Clean configuration and initialize all disks.
    (5) Maintenance mode boot.
    (6) Update flash from backup config.
    (7) Install new software first.
    (8) Reboot node.
    (9) Configure Advanced Drive Partitioning.
    Selection (1-9)?  wipeconfig
This option deletes critical system configuration, including cluster
membership.
Warning: do not run this option on a HA node that has been taken
over.
Are you sure you want to continue?: yes
Rebooting to finish wipeconfig request.
```

Manually assigning drives to pool 0

If you did not receive the systems pre-configured from the factory, you might have to manually assign the pool 0 drives. Depending on the platform model and whether the system is using ADP, you must manually assign drives to pool 0 for each node in the MetroCluster IP configuration. The procedure you use depends on the version of ONTAP you are using.

Manually assigning drives for pool 0 (ONTAP 9.4 and later)

If the system has not been pre-configured at the factory and does not meet the requirements for automatic drive assignment, you must manually assign the pool 0 drives.

About this task

This procedure applies to configurations running ONTAP 9.4 or later.

To determine if your system requires manual disk assignment, you should review Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later.

You perform these steps in Maintenance mode. The procedure must be performed on each node in the configuration.

Examples in this section are based on the following assumptions:

- node_A_1 and node_A_2 own drives on:

◦ site_A-shelf_1 (local)

  ◦ site_B-shelf_2 (remote)

- node_B_1 and node_B_2 own drives on:

  ◦ site_B-shelf_1 (local)

  ◦ site_A-shelf_2 (remote)

Steps

1. Display the boot menu:

    boot_ontap menu

2. Select option 9a.

    The following screen shows the boot menu prompt:

    ```
    Please choose one of the following:

        (1) Normal Boot.
        (2) Boot without /etc/rc.
        (3) Change password.
        (4) Clean configuration and initialize all disks.
        (5) Maintenance mode boot.
        (6) Update flash from backup config.
        (7) Install new software first.
        (8) Reboot node.
        (9) Configure Advanced Drive Partitioning.
        Selection (1-9)?  9a
    ########## WARNING ##########

        This is a disruptive operation and will result in the
        loss of all filesystem data. Before proceeding further,
        make sure that:
        1) This option (9a) has been executed or will be executed
        on the HA partner node (and DR/DR-AUX partner nodes if
        applicable), prior to reinitializing any system in the
        HA-pair (or MetroCluster setup).
        2) The HA partner node (and DR/DR-AUX partner nodes if
        applicable) is currently waiting at the boot menu.

        Do you still want to continue (yes/no)? yes
    ```

3. When the node restarts, press Ctrl-C when prompted to display the boot menu and then select the option for **Maintenance mode boot**.

4. In Maintenance mode, manually assign drives for the local aggregates on the node:

    disk assign *disk-id* -p 0 -s *local-node-sysid*

The drives should be assigned symmetrically, so each node has an equal number of drives. The following steps are for a configuration with two storage shelves at each site.

    a. When configuring node_A_1, manually assign drives from slot 0 to 11 to pool0 of node A1 from site_A-shelf_1.

    b. When configuring node_A_2, manually assign drives from slot 12 to 23 to pool0 of node A2 from site_A-shelf_1.

    c. When configuring node_B_1, manually assign drives from slot 0 to 11 to pool0 of node B1 from site_B-shelf_1.

    d. When configuring node_B_2, manually assign drives from slot 12 to 23 to pool0 of node B2 from site_B-shelf_1.

5. Exit Maintenance mode:

```
halt
```

6. Display the boot menu:

```
boot_ontap menu
```

7. Select option "4" from the boot menu and let the system boot.

8. Repeat these steps on the other nodes in the MetroCluster IP configuration.

9. Proceed to Setting up ONTAP.

Manually assigning drives for pool 0 (ONTAP 9.3)

If you have at least two disk shelves for each node, you use ONTAP's auto-assignment functionality to automatically assign the local (pool 0) disks.

About this task

While the node is in Maintenance mode, you must first assign a single disk on the appropriate shelves to pool 0. ONTAP then automatically assigns the rest of the disks on the shelf to the same pool. This task is not required on systems received from the factory, which have pool 0 to contain the pre-configured root aggregate.

This procedure applies to configurations running ONTAP 9.3.

This procedure is not required if you received your MetroCluster configuration from the factory. Nodes from the factory are configured with pool 0 disks and root aggregates.

This procedure can be used only if you have at least two disk shelves for each node, which allows shelf-level autoassignment of disks. If you cannot use shelf-level autoassignment, you must manually assign your local disks so that each node has a local pool of disks (pool 0).

These steps must be performed in Maintenance mode.

Examples in this section assume the following disk shelves:

- node_A_1 owns disks on:
  - site_A-shelf_1 (local)
  - site_B-shelf_2 (remote)
- node_A_2 is connected to:

- site_A-shelf_3 (local)
- site_B-shelf_4 (remote)
- node_B_1 is connected to:
  - site_B-shelf_1 (local)
  - site_A-shelf_2 (remote)
- node_B_2 is connected to:
  - site_B-shelf_3 (local)
  - site_A-shelf_4 (remote)

Steps

1. Manually assign a single disk for root aggregate on each node:

   ```
   disk assign disk-id -p 0 -s local-node-sysid
   ```

   The manual assignment of these disks allows the ONTAP autoassignment feature to assign the rest of the disks on each shelf.

   a. On node_A_1, manually assign one disk from local site_A-shelf_1 to pool 0.

   b. On node_A_2, manually assign one disk from local site_A-shelf_3 to pool 0.

   c. On node_B_1, manually assign one disk from local site_B-shelf_1 to pool 0.

   d. On node_B_2, manually assign one disk from local site_B-shelf_3 to pool 0.

2. Boot each node at site A, using option 4 on the boot menu:

   You should complete this step on a node before proceeding to the next node.

   a. Exit Maintenance mode:

   ```
   halt
   ```

   b. Display the boot menu:

   ```
   boot_ontap menu
   ```

   c. Select option 4 from the boot menu and proceed.

3. Boot each node at site B, using option 4 on the boot menu:

   You should complete this step on a node before proceeding to the next node.

   a. Exit Maintenance mode:

   ```
   halt
   ```

   b. Display the boot menu:

   ```
   boot_ontap menu
   ```

   c. Select option 4 from the boot menu and proceed.

Setting up ONTAP

After you boot each node, you are prompted to perform basic node and cluster configuration. After configuring the cluster, you return to the ONTAP CLI to create aggregates and create the MetroCluster configuration.

Before you begin

- You must have cabled the MetroCluster configuration.
- You must not have configured the Service Processor.

If you need to netboot the new controllers, see Netbooting the new controller modules.

About this task

This task must be performed on both clusters in the MetroCluster configuration.

Steps

1. Power up each node at the local site if you have not already done so and let them all boot completely.

   If the system is in Maintenance mode, you need to issue the halt command to exit Maintenance mode, and then issue the `boot_ontap` command to boot the system and get to cluster setup.

2. On the first node in each cluster, proceed through the prompts to configure the cluster.

   a. Enable the AutoSupport tool by following the directions provided by the system.

   The output should be similar to the following:

```
Welcome to the cluster setup wizard.

    You can enter the following commands at any time:
    "help" or "?" - if you want to have a question clarified,
    "back" - if you want to change previously answered questions, and
    "exit" or "quit" - if you want to quit the cluster setup wizard.
    Any changes you made before quitting will be saved.

    You can return to cluster setup at any time by typing "cluster
setup".
    To accept a default or omit a question, do not enter a value.

    This system will send event messages and periodic reports to
NetApp Technical
    Support. To disable this feature, enter
    autosupport modify -support disable
    within 24 hours.

    Enabling AutoSupport can significantly speed problem
determination and
    resolution should a problem occur on your system.
    For further information on AutoSupport, see:
    http://support.netapp.com/autosupport/

    Type yes to confirm and continue {yes}: yes
.
.
.
```

b. Configure the node management interface by responding to the prompts.

The prompts are similar to the following:

```
Enter the node management interface port [e0M]:
Enter the node management interface IP address: 172.17.8.229
Enter the node management interface netmask: 255.255.254.0
Enter the node management interface default gateway: 172.17.8.1
A node management interface on port e0M with IP address 172.17.8.229
has been created.
```

c. Create the cluster by responding to the prompts.

The prompts are similar to the following:

```
Do you want to create a new cluster or join an existing cluster?
{create, join}:
create


Do you intend for this node to be used as a single node cluster?
{yes, no} [no]:
no

Existing cluster interface configuration found:

Port MTU IP Netmask
e0a 1500 169.254.18.124 255.255.0.0
e1a 1500 169.254.184.44 255.255.0.0

Do you want to use this configuration? {yes, no} [yes]: no

System Defaults:
Private cluster network ports [e0a,e1a].
Cluster port MTU values will be set to 9000.
Cluster interface IP addresses will be automatically generated.

Do you want to use these defaults? {yes, no} [yes]: no

Enter the cluster administrator's (username "admin") password:

Retype the password:


Step 1 of 5: Create a Cluster
You can type "back", "exit", or "help" at any question.

List the private cluster network ports [e0a,e1a]:
Enter the cluster ports' MTU size [9000]:
Enter the cluster network netmask [255.255.0.0]: 255.255.254.0
Enter the cluster interface IP address for port e0a: 172.17.10.228
Enter the cluster interface IP address for port e1a: 172.17.10.229
Enter the cluster name: cluster_A

Creating cluster cluster_A

Starting cluster support services ...

Cluster cluster_A has been created.
```

d. Add licenses, set up a Cluster Administration SVM, and enter DNS information by responding to the prompts.

The prompts are similar to the following:

```
Step 2 of 5: Add Feature License Keys
You can type "back", "exit", or "help" at any question.


Enter an additional license key []:



Step 3 of 5: Set Up a Vserver for Cluster Administration
You can type "back", "exit", or "help" at any question.



Enter the cluster management interface port [e3a]:
Enter the cluster management interface IP address: 172.17.12.153
Enter the cluster management interface netmask: 255.255.252.0
Enter the cluster management interface default gateway: 172.17.12.1

A cluster management interface on port e3a with IP address
172.17.12.153 has been created. You can use this address to connect
to and manage the cluster.

Enter the DNS domain names: lab.netapp.com
Enter the name server IP addresses: 172.19.2.30
DNS lookup for the admin Vserver will use the lab.netapp.com domain.

Step 4 of 5: Configure Storage Failover (SFO)
You can type "back", "exit", or "help" at any question.



SFO will be enabled when the partner joins the cluster.



Step 5 of 5: Set Up the Node
You can type "back", "exit", or "help" at any question.

Where is the controller located []: svl
```

e. Enable storage failover and set up the node by responding to the prompts.

The prompts are similar to the following:

```
Step 4 of 5: Configure Storage Failover (SFO)
You can type "back", "exit", or "help" at any question.


SFO will be enabled when the partner joins the cluster.


Step 5 of 5: Set Up the Node
You can type "back", "exit", or "help" at any question.

Where is the controller located []: site_A
```

    f. Complete the configuration of the node, but do not create data aggregates.

    You can use ONTAP System Manager, pointing your web browser to the cluster management IP address (https://172.17.12.153).

    Cluster management using System Manager (Versions 9.0 to 9.6)

    ONTAP System Manager (Version 9.7 and later)

3. Boot the next controller and join it to the cluster, following the prompts.

4. Confirm that nodes are configured in high-availability mode:

```
storage failover show -fields mode
```

If not, you must configure HA mode on each node, and then reboot the nodes:

```
storage failover modify -mode ha -node localhost
```

This command configures high-availability mode but does not enable storage failover. Storage failover is automatically enabled when you configure the MetroCluster configuration later in the process.

5. Confirm that you have four ports configured as cluster interconnects:

```
network port show
```

The MetroCluster IP interfaces are not configured at this time and do not appear in the command output.

The following example shows two cluster ports on node_A_1:

```
cluster_A::*> network port show -role cluster



Node: node_A_1


Ignore
```

```
                                                   Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status

--------- ------------ ---------------- ---- ---- ----------- --------
------

e4a       Cluster      Cluster          up   9000 auto/40000 healthy
false

e4e       Cluster      Cluster          up   9000 auto/40000 healthy
false


Node: node_A_2


Ignore

                                                   Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status

--------- ------------ ---------------- ---- ---- ----------- --------
------

e4a       Cluster      Cluster          up   9000 auto/40000 healthy
false

e4e       Cluster      Cluster          up   9000 auto/40000 healthy
false


4 entries were displayed.
```

6. Repeat these steps on the partner cluster.

What to do next

Return to the ONTAP command-line interface and complete the MetroCluster configuration by performing the tasks that follow.

Configuring the clusters into a MetroCluster configuration

You must peer the clusters, mirror the root aggregates, create a mirrored data aggregate, and then issue the command to implement the MetroCluster operations.

Disabling automatic drive assignment (if doing manual assignment in ONTAP 9.4)

In ONTAP 9.4, if your MetroCluster IP configuration has fewer than four external storage shelves per site, you must disable automatic drive assignment on all nodes and manually assign drives.

About this task

This task is not required in ONTAP 9.5 and later.

This task does not apply to an AFF A800 system with an internal shelf and no external shelves.

Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later

Steps

1. Disable automatic drive assignment:

   ```
   storage disk option modify -node node_name -autoassign off
   ```

2. You need to issue this command on all nodes in the MetroCluster IP configuration.

Verifying drive assignment of pool 0 drives

You must verify that the remote drives are visible to the nodes and have been assigned correctly.

About this task

Automatic assignment depends on the storage system platform model and drive shelf arrangement.

Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later

Steps

1. Verify that pool 0 drives are assigned automatically:

   ```
   disk show
   ```

   The following example shows the "cluster_A" output for an AFF A800 system with no external shelves.

   One quarter (8 drives) were automatically assigned to "node_A_1" and one quarter were automatically assigned to "node_A_2". The remaining drives will be remote (pool 1) drives for "node_B_1" and "node_B_2".

   ```
   cluster_A::*> disk show
                   Usable      Disk      Container           Container
   Disk            Size        Shelf Bay Type      Type      Name
   Owner
   --------------- ---------- ----- --- ------- ----------- ---------
   --------
   node_A_1:0n.12  1.75TB      0     12  SSD-NVM shared      aggr0
   ```

```
node_A_1
node_A_1:0n.13    1.75TB    0    13  SSD-NVM shared        aggr0
node_A_1
node_A_1:0n.14    1.75TB    0    14  SSD-NVM shared        aggr0
node_A_1
node_A_1:0n.15    1.75TB    0    15  SSD-NVM shared        aggr0
node_A_1
node_A_1:0n.16    1.75TB    0    16  SSD-NVM shared        aggr0
node_A_1
node_A_1:0n.17    1.75TB    0    17  SSD-NVM shared        aggr0
node_A_1
node_A_1:0n.18    1.75TB    0    18  SSD-NVM shared        aggr0
node_A_1
node_A_1:0n.19    1.75TB    0    19  SSD-NVM shared        -
node_A_1
node_A_2:0n.0     1.75TB    0    0   SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.1     1.75TB    0    1   SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.2     1.75TB    0    2   SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.3     1.75TB    0    3   SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.4     1.75TB    0    4   SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.5     1.75TB    0    5   SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.6     1.75TB    0    6   SSD-NVM shared
aggr0_node_A_2_0 node_A_2
node_A_2:0n.7     1.75TB    0    7   SSD-NVM shared        -
node_A_2
node_A_2:0n.24    -         0    24  SSD-NVM unassigned -        -
node_A_2:0n.25    -         0    25  SSD-NVM unassigned -        -
node_A_2:0n.26    -         0    26  SSD-NVM unassigned -        -
node_A_2:0n.27    -         0    27  SSD-NVM unassigned -        -
node_A_2:0n.28    -         0    28  SSD-NVM unassigned -        -
node_A_2:0n.29    -         0    29  SSD-NVM unassigned -        -
node_A_2:0n.30    -         0    30  SSD-NVM unassigned -        -
node_A_2:0n.31    -         0    31  SSD-NVM unassigned -        -
node_A_2:0n.36    -         0    36  SSD-NVM unassigned -        -
node_A_2:0n.37    -         0    37  SSD-NVM unassigned -        -
node_A_2:0n.38    -         0    38  SSD-NVM unassigned -        -
node_A_2:0n.39    -         0    39  SSD-NVM unassigned -        -
node_A_2:0n.40    -         0    40  SSD-NVM unassigned -        -
node_A_2:0n.41    -         0    41  SSD-NVM unassigned -        -
node_A_2:0n.42    -         0    42  SSD-NVM unassigned -        -
```

```
node_A_2:0n.43    -          0     43   SSD-NVM unassigned  -         -
32 entries were displayed.
```

The following example shows the "cluster_B" output:

```
cluster_B::> disk show
                   Usable    Disk              Container  Container
Disk               Size      Shelf Bay Type    Type       Name      Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------

Info: This cluster has partitioned disks. To get a complete list of
spare disk
capacity use "storage aggregate show-spare-disks".
node_B_1:0n.12   1.75TB     0     12  SSD-NVM shared      aggr0
node_B_1
node_B_1:0n.13   1.75TB     0     13  SSD-NVM shared      aggr0
node_B_1
node_B_1:0n.14   1.75TB     0     14  SSD-NVM shared      aggr0
node_B_1
node_B_1:0n.15   1.75TB     0     15  SSD-NVM shared      aggr0
node_B_1
node_B_1:0n.16   1.75TB     0     16  SSD-NVM shared      aggr0
node_B_1
node_B_1:0n.17   1.75TB     0     17  SSD-NVM shared      aggr0
node_B_1
node_B_1:0n.18   1.75TB     0     18  SSD-NVM shared      aggr0
node_B_1
node_B_1:0n.19   1.75TB     0     19  SSD-NVM shared      -
node_B_1
node_B_2:0n.0    1.75TB     0     0   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.1    1.75TB     0     1   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.2    1.75TB     0     2   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.3    1.75TB     0     3   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.4    1.75TB     0     4   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.5    1.75TB     0     5   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
node_B_2:0n.6    1.75TB     0     6   SSD-NVM shared
aggr0_node_B_1_0 node_B_2
```

```
   node_B_2:0n.7     1.75TB      0      7    SSD-NVM shared       -
   node_B_2
   node_B_2:0n.24     -          0     24    SSD-NVM unassigned   -        -
   node_B_2:0n.25     -          0     25    SSD-NVM unassigned   -        -
   node_B_2:0n.26     -          0     26    SSD-NVM unassigned   -        -
   node_B_2:0n.27     -          0     27    SSD-NVM unassigned   -        -
   node_B_2:0n.28     -          0     28    SSD-NVM unassigned   -        -
   node_B_2:0n.29     -          0     29    SSD-NVM unassigned   -        -
   node_B_2:0n.30     -          0     30    SSD-NVM unassigned   -        -
   node_B_2:0n.31     -          0     31    SSD-NVM unassigned   -        -
   node_B_2:0n.36     -          0     36    SSD-NVM unassigned   -        -
   node_B_2:0n.37     -          0     37    SSD-NVM unassigned   -        -
   node_B_2:0n.38     -          0     38    SSD-NVM unassigned   -        -
   node_B_2:0n.39     -          0     39    SSD-NVM unassigned   -        -
   node_B_2:0n.40     -          0     40    SSD-NVM unassigned   -        -
   node_B_2:0n.41     -          0     41    SSD-NVM unassigned   -        -
   node_B_2:0n.42     -          0     42    SSD-NVM unassigned   -        -
   node_B_2:0n.43     -          0     43    SSD-NVM unassigned   -        -
   32 entries were displayed.


   cluster_B::>
```

Peering the clusters

The clusters in the MetroCluster configuration must be in a peer relationship so that they can communicate with each other and perform the data mirroring essential to MetroCluster disaster recovery.

Related information

Cluster and SVM peering express configuration

Considerations when using dedicated ports

Considerations when sharing data ports

Configuring intercluster LIFs for cluster peering

You must create intercluster LIFs on ports used for communication between the MetroCluster partner clusters. You can use dedicated ports or ports that also have data traffic.

Configuring intercluster LIFs on dedicated ports

You can configure intercluster LIFs on dedicated ports. Doing so typically increases the available bandwidth for replication traffic.

Steps

1. List the ports in the cluster:

   ```
   network port show
   ```

For complete command syntax, see the man page.

The following example shows the network ports in "cluster01":

```
cluster01::> network port show
                                                           Speed
(Mbps)
Node    Port      IPspace       Broadcast Domain Link   MTU    Admin/Oper
------  --------- ------------  ---------------- -----  -------
------------
cluster01-01
        e0a       Cluster       Cluster          up     1500   auto/1000
        e0b       Cluster       Cluster          up     1500   auto/1000
        e0c       Default       Default          up     1500   auto/1000
        e0d       Default       Default          up     1500   auto/1000
        e0e       Default       Default          up     1500   auto/1000
        e0f       Default       Default          up     1500   auto/1000
cluster01-02
        e0a       Cluster       Cluster          up     1500   auto/1000
        e0b       Cluster       Cluster          up     1500   auto/1000
        e0c       Default       Default          up     1500   auto/1000
        e0d       Default       Default          up     1500   auto/1000
        e0e       Default       Default          up     1500   auto/1000
        e0f       Default       Default          up     1500   auto/1000
```

2. Determine which ports are available to dedicate to intercluster communication:

```
network interface show -fields home-port,curr-port
```

For complete command syntax, see the man page.

The following example shows that ports "e0e" and "e0f" have not been assigned LIFs:

```
cluster01::> network interface show -fields home-port,curr-port
vserver lif                      home-port curr-port
------- -------------------- --------- ---------
Cluster cluster01-01_clus1    e0a       e0a
Cluster cluster01-01_clus2    e0b       e0b
Cluster cluster01-02_clus1    e0a       e0a
Cluster cluster01-02_clus2    e0b       e0b
cluster01
        cluster_mgmt          e0c       e0c
cluster01
        cluster01-01_mgmt1    e0c       e0c
cluster01
        cluster01-02_mgmt1    e0c       e0c
```

3. Create a failover group for the dedicated ports:

```
network interface failover-groups create -vserver system_SVM -failover-group
failover_group -targets physical_or_logical_ports
```

The following example assigns ports "e0e" and" e0f" to failover group "intercluster01" on system "SVMcluster01":

```
cluster01::> network interface failover-groups create -vserver cluster01
-failover-group
intercluster01 -targets
cluster01-01:e0e,cluster01-01:e0f,cluster01-02:e0e,cluster01-02:e0f
```

4. Verify that the failover group was created:

```
network interface failover-groups show
```

For complete command syntax, see the man page.

```
cluster01::> network interface failover-groups show
                                    Failover
Vserver            Group            Targets
---------------- ----------------
-----------------------------------------------
Cluster
                    Cluster
                                    cluster01-01:e0a, cluster01-01:e0b,
                                    cluster01-02:e0a, cluster01-02:e0b
cluster01
                    Default
                                    cluster01-01:e0c, cluster01-01:e0d,
                                    cluster01-02:e0c, cluster01-02:e0d,
                                    cluster01-01:e0e, cluster01-01:e0f
                                    cluster01-02:e0e, cluster01-02:e0f
                    intercluster01
                                    cluster01-01:e0e, cluster01-01:e0f
                                    cluster01-02:e0e, cluster01-02:e0f
```

5. Create intercluster LIFs on the system SVM and assign them to the failover group.

ONTAP version	Command

9.6 and later	`network interface create -vserver system_SVM -lif LIF_name -service -policy default-intercluster -home -node node -home-port port -address port_IP -netmask netmask -failover -group failover_group`
9.5 and earlier	`network interface create -vserver system_SVM -lif LIF_name -role intercluster -home-node node -home -port port -address port_IP -netmask netmask -failover-group failover_group`

For complete command syntax, see the man page.

The following example creates intercluster LIFs "cluster01_icl01" and "cluster01_icl02" in failover group "intercluster01":

```
cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl01 -service-
policy default-intercluster -home-node cluster01-01 -home-port e0e
-address 192.168.1.201
-netmask 255.255.255.0 -failover-group intercluster01

cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl02 -service-
policy default-intercluster -home-node cluster01-02 -home-port e0e
-address 192.168.1.202
-netmask 255.255.255.0 -failover-group intercluster01
```

6. Verify that the intercluster LIFs were created:

In ONTAP 9.6 and later:
`network interface show -service-policy default-intercluster`
In ONTAP 9.5 and earlier:
`network interface show -role intercluster`

For complete command syntax, see the man page.

```
cluster01::> network interface show -service-policy default-intercluster
            Logical     Status     Network                 Current
Current Is
Vserver      Interface  Admin/Oper Address/Mask            Node          Port
Home
----------- ---------- ---------- ------------------ -------------
------- ----
cluster01
            cluster01_icl01
                          up/up       192.168.1.201/24   cluster01-01  e0e
true
            cluster01_icl02
                          up/up       192.168.1.202/24   cluster01-02  e0f
true
```

7. Verify that the intercluster LIFs are redundant:

In ONTAP 9.6 and later:
`network interface show -service-policy default-intercluster -failover`

In ONTAP 9.5 and earlier:
`network interface show -role intercluster -failover`

For complete command syntax, see the man page.

The following example shows that the intercluster LIFs "cluster01_icl01", and "cluster01_icl02" on the "SVMe0e" port will fail over to the "e0f" port.

```
cluster01::> network interface show -service-policy default-intercluster
-failover
         Logical          Home                      Failover        Failover
Vserver  Interface        Node:Port                 Policy          Group
-------- --------------- --------------------- --------------- --------
cluster01
         cluster01_icl01 cluster01-01:e0e   local-only
intercluster01
                             Failover Targets:  cluster01-01:e0e,
                                                cluster01-01:e0f
         cluster01_icl02 cluster01-02:e0e   local-only
intercluster01
                             Failover Targets:  cluster01-02:e0e,
                                                cluster01-02:e0f
```

Related information

Configuring intercluster LIFs on shared data ports

You can configure intercluster LIFs on ports shared with the data network. Doing so reduces the number of ports you need for intercluster networking.

Steps

1. List the ports in the cluster:

   network port show

   For complete command syntax, see the man page.

   The following example shows the network ports in "cluster01":

   ```
   cluster01::> network port show
                                                       Speed
    (Mbps)
    Node    Port      IPspace       Broadcast Domain Link   MTU    Admin/Oper
    ------  --------- ------------  ---------------- -----  -------
    ------------
    cluster01-01
            e0a       Cluster       Cluster          up     1500   auto/1000
            e0b       Cluster       Cluster          up     1500   auto/1000
            e0c       Default       Default          up     1500   auto/1000
            e0d       Default       Default          up     1500   auto/1000
    cluster01-02
            e0a       Cluster       Cluster          up     1500   auto/1000
            e0b       Cluster       Cluster          up     1500   auto/1000
            e0c       Default       Default          up     1500   auto/1000
            e0d       Default       Default          up     1500   auto/1000
   ```

2. Create intercluster LIFs on the system SVM:

In ONTAP 9.6 and later:
network interface create -vserver *system_SVM* -lif *LIF_name* -service-policy default-intercluster -home-node *node* -home-port *port* -address *port_IP* -netmask *netmask*
In ONTAP 9.5 and earlier:
network interface create -vserver *system_SVM* -lif *LIF_name* -role intercluster -home-node *node* -home-port *port* -address *port_IP* -netmask *netmask*

   For complete command syntax, see the man page.

The following example creates intercluster LIFs "cluster01_icl01" and "cluster01_icl02":

```
cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl01 -service-
policy default-intercluster -home-node cluster01-01 -home-port e0c
-address 192.168.1.201
-netmask 255.255.255.0

cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl02 -service-
policy default-intercluster -home-node cluster01-02 -home-port e0c
-address 192.168.1.202
-netmask 255.255.255.0
```

3. Verify that the intercluster LIFs were created:

In ONTAP 9.6 and later:
`network interface show -service-policy default-intercluster`
In ONTAP 9.5 and earlier:
`network interface show -role intercluster`

For complete command syntax, see the man page.

```
cluster01::> network interface show -service-policy default-intercluster
            Logical    Status     Network                 Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask         Node          Port
Home
----------- ---------- ---------- ------------------ -------------
------- ----
cluster01
            cluster01_icl01
                       up/up      192.168.1.201/24   cluster01-01  e0c
true
            cluster01_icl02
                       up/up      192.168.1.202/24   cluster01-02  e0c
true
```

4. Verify that the intercluster LIFs are redundant:

In ONTAP 9.6 and later:

```
network interface show -service-policy default-intercluster -failover
```

In ONTAP 9.5 and earlier:

```
network interface show -role intercluster -failover
```

For complete command syntax, see the man page.

The following example shows that intercluster LIFs "cluster01_icl01" and "cluster01_icl02" on the "e0c" port will fail over to the "e0d" port.

```
cluster01::> network interface show -service-policy default-intercluster
-failover
          Logical          Home                       Failover          Failover
Vserver   Interface        Node:Port                  Policy            Group
--------  ---------------  --------------------  ---------------  --------
cluster01
          cluster01_icl01 cluster01-01:e0c    local-only
192.168.1.201/24
                              Failover Targets: cluster01-01:e0c,
                                                cluster01-01:e0d
          cluster01_icl02 cluster01-02:e0c    local-only
192.168.1.201/24
                              Failover Targets: cluster01-02:e0c,
                                                cluster01-02:e0d
```

Related information

Considerations when sharing data ports

Creating a cluster peer relationship

You can use the cluster peer create command to create a peer relationship between a local and remote cluster. After the peer relationship has been created, you can run cluster peer create on the remote cluster to authenticate it to the local cluster.

About this task

- You must have created intercluster LIFs on every node in the clusters that are being peered.

- The clusters must be running ONTAP 9.3 or later.

Steps

1. On the destination cluster, create a peer relationship with the source cluster:

   ```
   cluster peer create -generate-passphrase -offer-expiration MM/DD/YYYY
   HH:MM:SS|1…7days|1…168hours -peer-addrs peer_LIF_IPs -ipspace ipspace
   ```

   If you specify both -generate-passphrase and -peer-addrs, only the cluster whose intercluster LIFs are specified in -peer-addrs can use the generated password.

You can ignore the `-ipspace` option if you are not using a custom IPspace. For complete command syntax, see the man page.

The following example creates a cluster peer relationship on an unspecified remote cluster:

```
cluster02::> cluster peer create -generate-passphrase -offer-expiration
2days

                       Passphrase: UCa+6lRVICXeL/gq1WrK7ShR
               Expiration Time: 6/7/2017 08:16:10 EST
   Initial Allowed Vserver Peers: -
               Intercluster LIF IP: 192.140.112.101
                Peer Cluster Name: Clus_7ShR (temporary generated)

Warning: make a note of the passphrase - it cannot be displayed again.
```

2. On the source cluster, authenticate the source cluster to the destination cluster:

```
cluster peer create -peer-addrs peer_LIF_IPs -ipspace ipspace
```

For complete command syntax, see the man page.

The following example authenticates the local cluster to the remote cluster at intercluster LIF IP addresses "192.140.112.101" and "192.140.112.102":

```
cluster01::> cluster peer create -peer-addrs
192.140.112.101,192.140.112.102

Notice: Use a generated passphrase or choose a passphrase of 8 or more
characters.
        To ensure the authenticity of the peering relationship, use a
phrase or sequence of characters that would be hard to guess.

Enter the passphrase:
Confirm the passphrase:

Clusters cluster02 and cluster01 are peered.
```

Enter the passphrase for the peer relationship when prompted.

3. Verify that the cluster peer relationship was created:

```
cluster peer show -instance
```

```
cluster01::> cluster peer show -instance

                            Peer Cluster Name: cluster02
                Remote Intercluster Addresses: 192.140.112.101,
192.140.112.102
            Availability of the Remote Cluster: Available
                          Remote Cluster Name: cluster2
                          Active IP Addresses: 192.140.112.101,
192.140.112.102
                        Cluster Serial Number: 1-80-123456
                Address Family of Relationship: ipv4
            Authentication Status Administrative: no-authentication
              Authentication Status Operational: absent
                            Last Update Time: 02/05 21:05:41
                    IPspace for the Relationship: Default
```

4. Check the connectivity and status of the nodes in the peer relationship:

`cluster peer health show`

```
cluster01::> cluster peer health show
Node       cluster-Name                  Node-Name
          Ping-Status                   RDB-Health Cluster-Health  Avail…
---------- --------------------------- ---------   ---------------
--------
cluster01-01
          cluster02                     cluster02-01
            Data: interface_reachable
            ICMP: interface_reachable true        true            true
                                        cluster02-02
            Data: interface_reachable
            ICMP: interface_reachable true        true            true
cluster01-02
          cluster02                     cluster02-01
            Data: interface_reachable
            ICMP: interface_reachable true        true            true
                                        cluster02-02
            Data: interface_reachable
            ICMP: interface_reachable true        true            true
```

Creating the DR group

You must create the disaster recovery (DR) group relationships between the clusters.

About this task

You perform this procedure on one of the clusters in the MetroCluster configuration to create the DR relationships between the nodes in both clusters.

ⓘ The DR relationships cannot be changed after the DR groups are created.

Steps

1. Verify that the nodes are ready for creation of the DR group by entering the following command on each node:

   `metrocluster configuration-settings show-status`

   The command output should show that the nodes are ready:

   ```
   cluster_A::> metrocluster configuration-settings show-status
   Cluster                       Node          Configuration Settings Status
   --------------------------- -------------
   -------------------------------
   cluster_A                     node_A_1      ready for DR group create
                                 node_A_2      ready for DR group create
   2 entries were displayed.
   ```

   ```
   cluster_B::> metrocluster configuration-settings show-status
   Cluster                       Node          Configuration Settings Status
   --------------------------- -------------
   -------------------------------
   cluster_B                     node_B_1      ready for DR group create
                                 node_B_2      ready for DR group create
   2 entries were displayed.
   ```

2. Create the DR group:

```
metrocluster configuration-settings dr-group create -partner-cluster partner-
cluster-name -local-node local-node-name -remote-node remote-node-name
```

This command is issued only once. It does not need to be repeated on the partner cluster. In the command, you specify the name of the remote cluster and the name of one local node and one node on the partner cluster.

The two nodes you specify are configured as DR partners and the other two nodes (which are not specified in the command) are configured as the second DR pair in the DR group. These relationships cannot be changed after you enter this command.

The following command creates these DR pairs:

- node_A_1 and node_B_1
- node_A_2 and node_B_2

```
Cluster_A::> metrocluster configuration-settings dr-group create
-partner-cluster cluster_B -local-node node_A_1 -remote-node node_B_1
[Job 27] Job succeeded: DR Group Create is successful.
```

Configuring and connecting the MetroCluster IP interfaces

You must configure the MetroCluster IP interfaces that are used for replication of each node's storage and nonvolatile cache. You then establish the connections using the MetroCluster IP interfaces. This creates iSCSI connections for storage replication.

About this task

> ℹ️ You must choose the MetroCluster IP addresses carefully because you cannot change them after initial configuration.

- You must create two interfaces for each node. The interfaces must be associated with the VLANs defined in the MetroCluster RCF file.
- You must create all MetroCluster IP interface "A" ports in the same VLAN and all MetroCluster IP interface "B" ports in the other VLAN. Refer to Considerations for MetroCluster IP configuration.

+

> ℹ️ - Certain platforms use a VLAN for the MetroCluster IP interface. By default, each of the two ports use a different VLAN: 10 and 20. You can also specify a different (non-default) VLAN higher than 100 (between 101 and 4095) using the `-vlan-id parameter` in the `metrocluster configuration-settings interface create` command.
> - Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also specify the `-gateway` parameter when creating MetroCluster IP interfaces. Refer to Considerations for layer 3 wide-area networks.

+ The following platform models use VLANs and allow configuration of a non-default VLAN ID.

+

AFF platforms	FAS platforms
• AFF A220 • AFF A250 • AFF A400	• FAS2750 • FAS500f • FAS8300 • FAS8700

The following IP addresses and subnets are used in the examples:

Node	Interface	IP address	Subnet
node_A_1	MetroCluster IP interface 1	10.1.1.1	10.1.1/24
	MetroCluster IP interface 2	10.1.2.1	10.1.2/24
node_A_2	MetroCluster IP interface 1	10.1.1.2	10.1.1/24
	MetroCluster IP interface 2	10.1.2.2	10.1.2/24
node_B_1	MetroCluster IP interface 1	10.1.1.3	10.1.1/24
	MetroCluster IP interface 2	10.1.2.3	10.1.2/24
node_B_2	MetroCluster IP interface 1	10.1.1.4	10.1.1/24
	MetroCluster IP interface 2	10.1.2.4	10.1.2/24

The physical ports used by the MetroCluster IP interfaces depends on the platform model, as shown in the following table.

Platform model	MetroCluster IP port	Note
AFF A900	e5b	
	e7b	

Platform model	MetroCluster IP port	Note
AFF A800	e0b	
	e1b	
AFF A700 and FAS900	e5a	
	e5b	
AFF A400	e1a	
	e1b	
AFF A320	e0g	
	e0h	
AFF A300 and FAS8200	e1a	
	e1b	
AFF A220 and FAS2750	e0a	On these systems, these physical ports are also used as cluster interfaces.
	e0b	
AFF A250 and FAS500f	e0c	
	e0d	
FAS8300 and FAS8700	e1a	
	e1b	

The port usage in the following examples is for an AFF A700 or a FAS9000 system.

Steps

1. Confirm that each node has disk automatic assignment enabled:

   ```
   storage disk option show
   ```

   Disk automatic assignment will assign pool 0 and pool 1 disks on a shelf-by-shelf basis.

   The Auto Assign column indicates whether disk automatic assignment is enabled.

```
Node          BKg. FW. Upd.  Auto Copy   Auto Assign  Auto Assign Policy
----------    -------------  ----------  -----------  ------------------
node_A_1                on          on           on             default
node_A_2                on          on           on             default
2 entries were displayed.
```

2. Verify you can create MetroCluster IP interfaces on the nodes:

   `metrocluster configuration-settings show-status`

   All nodes should be ready:

```
Cluster        Node          Configuration Settings Status
----------     -----------   --------------------------------
cluster_A
               node_A_1      ready for interface create
               node_A_2      ready for interface create
cluster_B
               node_B_1      ready for interface create
               node_B_2      ready for interface create
4 entries were displayed.
```

3. Create the interfaces on node_A_1.

   (i)
   - ◦ The port usage in the following examples is for an AFF A700 or a FAS9000 system (e5a and e5b). You must configure the interfaces on the correct ports for your platform model, as given above.
   - ◦ Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also specify the `-gateway` parameter when creating MetroCluster IP interfaces. Refer to Considerations for layer 3 wide-area networks.
   - ◦ On platform models that support VLANs for the MetroCluster IP interface, you can include the `-vlan-id` parameter if you don't want to use the default VLAN IDs.

   a. Configure the interface on port "e5a" on "node_A_1":

      `metrocluster configuration-settings interface create -cluster-name cluster-name -home-node node-name -home-port e5a -address ip-address -netmask netmask`

      The following example shows the creation of the interface on port "e5a" on "node_A_1" with IP address "10.1.1.1":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_1 -home-port e5a -address
10.1.1.1 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

b. Configure the interface on port "e5b" on "node_A_1":

```
metrocluster configuration-settings interface create -cluster-name cluster-
name -home-node node-name -home-port e5b -address ip-address -netmask
netmask
```

The following example shows the creation of the interface on port "e5b" on "node_A_1" with IP address
"10.1.2.1":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_1 -home-port e5b -address
10.1.2.1 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

(i) You can verify that these interfaces are present using the `metrocluster`
`configuration-settings interface show` command.

4. Create the interfaces on node_A_2.

(i)
  ◦ The port usage in the following examples is for an AFF A700 or a FAS9000 system (e5a
    and e5b). You must configure the interfaces on the correct ports for your platform model,
    as given above.

  ◦ Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also
    specify the `-gateway` parameter when creating MetroCluster IP interfaces. Refer to
    Considerations for layer 3 wide-area networks.

  ◦ On platform models that support VLANs for the MetroCluster IP interface, you can
    include the `-vlan-id` parameter if you don't want to use the default VLAN IDs.

a. Configure the interface on port "e5a" on "node_A_2":

```
metrocluster configuration-settings interface create -cluster-name cluster-
name -home-node node-name -home-port e5a -address ip-address -netmask
netmask
```

The following example shows the creation of the interface on port "e5a" on "node_A_2" with IP address
"10.1.1.2":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_2 -home-port e5a -address
10.1.1.2 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

On platform models that support VLANs for the MetroCluster IP interface, you can include the `-vlan` `-id` parameter if you don't want to use the default VLAN IDs. The following example shows the command for an AFF A220 system with a VLAN ID of 120:

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_2 -home-port e0a -address
10.1.1.2 -netmask 255.255.255.0 -vlan-id 120
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

b. Configure the interface on port "e5b" on "node_A_2":

```
metrocluster configuration-settings interface create -cluster-name cluster-
name -home-node node-name -home-port e5b -address ip-address -netmask
netmask
```

The following example shows the creation of the interface on port "e5b" on "node_A_2" with IP address "10.1.2.2":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_2 -home-port e5b -address
10.1.2.2 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

On platform models that support VLANs for the MetroCluster IP interface, you can include the `-vlan` `-id` parameter if you don't want to use the default VLAN IDs. The following example shows the command for an AFF A220 system with a VLAN ID of 220:

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_2 -home-port e0b -address
10.1.2.2 -netmask 255.255.255.0 -vlan-id 220
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

5. Create the interfaces on "node_B_1".

- The port usage in the following examples is for an AFF A700 or a FAS9000 system (e5a and e5b). You must configure the interfaces on the correct ports for your platform model, as given above.

- Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also specify the `-gateway` parameter when creating MetroCluster IP interfaces. Refer to Considerations for layer 3 wide-area networks.

- On platform models that support VLANs for the MetroCluster IP interface, you can include the `-vlan-id` parameter if you don't want to use the default VLAN IDs.

a. Configure the interface on port "e5a" on "node_B_1":

metrocluster configuration-settings interface create -cluster-name *cluster-name* -home-node *node-name* -home-port e5a -address *ip-address* -netmask *netmask*

The following example shows the creation of the interface on port "e5a" on "node_B_1" with IP address "10.1.1.3":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_1 -home-port e5a -address
10.1.1.3 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.cluster_B::>
```

b. Configure the interface on port "e5b" on "node_B_1":

metrocluster configuration-settings interface create -cluster-name *cluster-name* -home-node *node-name* -home-port e5a -address *ip-address* -netmask *netmask*

The following example shows the creation of the interface on port "e5b" on "node_B_1" with IP address "10.1.2.3":

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_1 -home-port e5b -address
10.1.2.3 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.cluster_B::>
```

6. Create the interfaces on "node_B_2".

- The port usage in the following examples is for an AFF A700 or a FAS9000 system (e5a and e5b). You must configure the interfaces on the correct ports for your platform model, as given above.

- Beginning with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also specify the `-gateway` parameter when creating MetroCluster IP interfaces. Refer to Considerations for layer 3 wide-area networks.

- On platform models that support VLANs for the MetroCluster IP interface, you can include the `-vlan-id` parameter if you don't want to use the default VLAN IDs.

a. Configure the interface on port e5a on node_B_2:

```
metrocluster configuration-settings interface create -cluster-name cluster-name -home-node node-name -home-port e5a -address ip-address -netmask netmask
```

The following example shows the creation of the interface on port "e5a" on "node_B_2" with IP address "10.1.1.4":

```
cluster_B::>metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_2 -home-port e5a -address
10.1.1.4 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.cluster_A::>
```

b. Configure the interface on port "e5b" on "node_B_2":

```
metrocluster configuration-settings interface create -cluster-name cluster-name -home-node node-name -home-port e5b -address ip-address -netmask netmask
```

The following example shows the creation of the interface on port "e5b" on "node_B_2" with IP address "10.1.2.4":

```
cluster_B::> metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_2 -home-port e5b -address
10.1.2.4 -netmask 255.255.255.0
[Job 28] Job succeeded: Interface Create is successful.
cluster_A::>
```

7. Verify that the interfaces have been configured:

```
metrocluster configuration-settings interface show
```

The following example shows that the configuration state for each interface is completed.

```
cluster_A::> metrocluster configuration-settings interface show
DR                                                            Config
Group Cluster Node    Network Address Netmask         Gateway  State
----- ------- ------- --------------- --------------- ---------
----------
1     cluster_A  node_A_1
                Home Port: e5a
                       10.1.1.1     255.255.255.0   -         completed
                Home Port: e5b
                       10.1.2.1     255.255.255.0   -         completed
                node_A_2
                Home Port: e5a
                       10.1.1.2     255.255.255.0   -         completed
                Home Port: e5b
                       10.1.2.2     255.255.255.0   -         completed
      cluster_B  node_B_1
                Home Port: e5a
                       10.1.1.3     255.255.255.0   -         completed
                Home Port: e5b
                       10.1.2.3     255.255.255.0   -         completed
                node_B_2
                Home Port: e5a
                       10.1.1.4     255.255.255.0   -         completed
                Home Port: e5b
                       10.1.2.4     255.255.255.0   -         completed
8 entries were displayed.
cluster_A::>
```

8. Verify that the nodes are ready to connect the MetroCluster interfaces:

`metrocluster configuration-settings show-status`

The following example shows all nodes in the "ready for connection" state:

```
Cluster       Node         Configuration Settings Status
----------    -----------  ---------------------------------
cluster_A
              node_A_1     ready for connection connect
              node_A_2     ready for connection connect
cluster_B
              node_B_1     ready for connection connect
              node_B_2     ready for connection connect
4 entries were displayed.
```

9. Establish the connections: `metrocluster configuration-settings connection connect`

   The IP addresses cannot be changed after you issue this command.

   The following example shows cluster_A is successfully connected:

   ```
   cluster_A::> metrocluster configuration-settings connection connect
   [Job 53] Job succeeded: Connect is successful.
   cluster_A::>
   ```

10. Verify that the connections have been established:

    ```
    metrocluster configuration-settings show-status
    ```

    The configuration settings status for all nodes should be completed:

    ```
    Cluster         Node        Configuration Settings Status
    ----------      ----------  --------------------------------
    cluster_A
                    node_A_1    completed
                    node_A_2    completed
    cluster_B
                    node_B_1    completed
                    node_B_2    completed
    4 entries were displayed.
    ```

11. Verify that the iSCSI connections have been established:

    a. Change to the advanced privilege level:

    ```
    set -privilege advanced
    ```

    You need to respond with `y` when you are prompted to continue into advanced mode and you see the advanced mode prompt (`*>`).

    b. Display the connections:

    ```
    storage iscsi-initiator show
    ```

    On systems running ONTAP 9.5, there are eight MetroCluster IP initiators on each cluster that should appear in the output.

    On systems running ONTAP 9.4 and earlier, there are four MetroCluster IP initiators on each cluster that should appear in the output.

    The following example shows the eight MetroCluster IP initiators on a cluster running ONTAP 9.5:

    ```
    cluster_A::*> storage iscsi-initiator show
    ```

```
Node Type Label    Target Portal            Target Name
Admin/Op
---- ---- -------- -----------------
------------------------------ --------


cluster_A-01
     dr_auxiliary
               mccip-aux-a-initiator
                    10.227.16.113:65200      prod506.com.company:abab44
up/up
               mccip-aux-a-initiator2
                    10.227.16.113:65200      prod507.com.company:abab44
up/up
               mccip-aux-b-initiator
                    10.227.95.166:65200      prod506.com.company:abab44
up/up
               mccip-aux-b-initiator2
                    10.227.95.166:65200      prod507.com.company:abab44
up/up
     dr_partner
               mccip-pri-a-initiator
                    10.227.16.112:65200      prod506.com.company:cdcd88
up/up
               mccip-pri-a-initiator2
                    10.227.16.112:65200      prod507.com.company:cdcd88
up/up
               mccip-pri-b-initiator
                    10.227.95.165:65200      prod506.com.company:cdcd88
up/up
               mccip-pri-b-initiator2
                    10.227.95.165:65200      prod507.com.company:cdcd88
up/up
cluster_A-02
     dr_auxiliary
               mccip-aux-a-initiator
                    10.227.16.112:65200      prod506.com.company:cdcd88
up/up
               mccip-aux-a-initiator2
                    10.227.16.112:65200      prod507.com.company:cdcd88
up/up
               mccip-aux-b-initiator
                    10.227.95.165:65200      prod506.com.company:cdcd88
up/up
               mccip-aux-b-initiator2
                    10.227.95.165:65200      prod507.com.company:cdcd88
up/up
```

```
        dr_partner
                mccip-pri-a-initiator
                        10.227.16.113:65200      prod506.com.company:abab44
    up/up
                mccip-pri-a-initiator2
                        10.227.16.113:65200      prod507.com.company:abab44
    up/up
                mccip-pri-b-initiator
                        10.227.95.166:65200      prod506.com.company:abab44
    up/up
                mccip-pri-b-initiator2
                        10.227.95.166:65200      prod507.com.company:abab44
    up/up
    16 entries were displayed.
```

   c. Return to the admin privilege level:

```
set -privilege admin
```

12. Verify that the nodes are ready for final implementation of the MetroCluster configuration:

```
metrocluster node show
```

```
cluster_A::> metrocluster node show
DR                                 Configuration  DR
Group Cluster Node                 State          Mirroring Mode
----- ------- ------------------   -------------- --------- ----
-       cluster_A
            node_A_1               ready to configure -     -
            node_A_2               ready to configure -     -
2 entries were displayed.
cluster_A::>
```

```
cluster_B::> metrocluster node show
DR                                 Configuration  DR
Group Cluster Node                 State          Mirroring Mode
----- ------- ------------------   -------------- --------- ----
-       cluster_B
            node_B_1               ready to configure -     -
            node_B_2               ready to configure -     -
2 entries were displayed.
cluster_B::>
```

Verifying or manually performing pool 1 drives assignment

Depending on the storage configuration, you must either verify pool 1 drive assignment or manually assign drives to pool 1 for each node in the MetroCluster IP configuration. The procedure you use depends on the version of ONTAP you are using.

Configuration type	Procedure
The systems meet the requirements for automatic drive assignment or, if running ONTAP 9.3, were received from the factory.	Verifying disk assignment for pool 1 disks
The configuration includes either three shelves, or, if it contains more than four shelves, has an uneven multiple of four shelves (for example, seven shelves), and is running ONTAP 9.5.	Manually assigning drives for pool 1 (ONTAP 9.4 or later)
The configuration does not include four storage shelves per site and is running ONTAP 9.4	Manually assigning drives for pool 1 (ONTAP 9.4 or later)
The systems were not received from the factory and are running ONTAP 9.3Systems received from the factory are pre-configured with assigned drives.	Manually assigning disks for pool 1 (ONTAP 9.3)

Verifying disk assignment for pool 1 disks

You must verify that the remote disks are visible to the nodes and have been assigned correctly.

Before you begin

You must wait at least ten minutes for disk auto-assignment to complete after the MetroCluster IP interfaces and connections were created with the `metrocluster configuration-settings connection connect` command.

Command output will show disk names in the form: node-name:0m.i1.0L1

Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later

Steps

1. Verify pool 1 disks are auto-assigned:

   `disk show`

   The following output shows the output for an AFF A800 system with no external shelves.

   Drive autoassignment has assigned one quarter (8 drives) to "node_A_1" and one quarter to "node_A_2". The remaining drives will be remote (pool 1) disks for "node_B_1" and "node_B_2".

   ```
   cluster_B::> disk show -host-adapter 0m -owner node_B_2
                        Usable      Disk                 Container   Container
   Disk                 Size        Shelf Bay Type       Type        Name
   Owner
   ```

```
----------------       ---------- ----- --- ------- ----------- ----------
--------
node_B_2:0m.i0.2L4   894.0GB    0      29  SSD-NVM shared       -
node_B_2
node_B_2:0m.i0.2L10 894.0GB    0      25  SSD-NVM shared       -
node_B_2
node_B_2:0m.i0.3L3   894.0GB    0      28  SSD-NVM shared       -
node_B_2
node_B_2:0m.i0.3L9   894.0GB    0      24  SSD-NVM shared       -
node_B_2
node_B_2:0m.i0.3L11 894.0GB    0      26  SSD-NVM shared       -
node_B_2
node_B_2:0m.i0.3L12 894.0GB    0      27  SSD-NVM shared       -
node_B_2
node_B_2:0m.i0.3L15 894.0GB    0      30  SSD-NVM shared       -
node_B_2
node_B_2:0m.i0.3L16 894.0GB    0      31  SSD-NVM shared       -
node_B_2
8 entries were displayed.

cluster_B::> disk show -host-adapter 0m -owner node_B_1
                   Usable     Disk                 Container   Container
Disk               Size       Shelf Bay Type       Type        Name
Owner
----------------       ---------- ----- --- ------- ----------- ----------
--------
node_B_1:0m.i2.3L19 1.75TB     0      42  SSD-NVM shared       -
node_B_1
node_B_1:0m.i2.3L20 1.75TB     0      43  SSD-NVM spare        Pool1
node_B_1
node_B_1:0m.i2.3L23 1.75TB     0      40  SSD-NVM shared       -
node_B_1
node_B_1:0m.i2.3L24 1.75TB     0      41  SSD-NVM spare        Pool1
node_B_1
node_B_1:0m.i2.3L29 1.75TB     0      36  SSD-NVM shared       -
node_B_1
node_B_1:0m.i2.3L30 1.75TB     0      37  SSD-NVM shared       -
node_B_1
node_B_1:0m.i2.3L31 1.75TB     0      38  SSD-NVM shared       -
node_B_1
node_B_1:0m.i2.3L32 1.75TB     0      39  SSD-NVM shared       -
node_B_1
8 entries were displayed.

cluster_B::> disk show
                   Usable     Disk                 Container   Container
```

```
Disk                  Size       Shelf Bay Type    Type        Name
Owner
---------------       ---------- ----- --- ------- ----------  ---------
--------
node_B_1:0m.i1.0L6  1.75TB       0     1    SSD-NVM shared      -
node_A_2
node_B_1:0m.i1.0L8  1.75TB       0     3    SSD-NVM shared      -
node_A_2
node_B_1:0m.i1.0L17 1.75TB       0     18  SSD-NVM shared       -
node_A_1
node_B_1:0m.i1.0L22 1.75TB       0     17 SSD-NVM shared - node_A_1
node_B_1:0m.i1.0L25 1.75TB       0     12 SSD-NVM shared - node_A_1
node_B_1:0m.i1.2L2  1.75TB       0     5 SSD-NVM shared - node_A_2
node_B_1:0m.i1.2L7  1.75TB       0     2 SSD-NVM shared - node_A_2
node_B_1:0m.i1.2L14 1.75TB       0     7 SSD-NVM shared - node_A_2
node_B_1:0m.i1.2L21 1.75TB       0     16 SSD-NVM shared - node_A_1
node_B_1:0m.i1.2L27 1.75TB       0     14 SSD-NVM shared - node_A_1
node_B_1:0m.i1.2L28 1.75TB       0     15 SSD-NVM shared - node_A_1
node_B_1:0m.i2.1L1  1.75TB       0     4 SSD-NVM shared - node_A_2
node_B_1:0m.i2.1L5  1.75TB       0     0 SSD-NVM shared - node_A_2
node_B_1:0m.i2.1L13 1.75TB       0     6 SSD-NVM shared - node_A_2
node_B_1:0m.i2.1L18 1.75TB       0     19 SSD-NVM shared - node_A_1
node_B_1:0m.i2.1L26 1.75TB       0     13 SSD-NVM shared - node_A_1
node_B_1:0m.i2.3L19 1.75TB       0 42 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L20 1.75TB       0 43 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L23 1.75TB       0 40 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L24 1.75TB       0 41 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L29 1.75TB       0 36 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L30 1.75TB       0 37 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L31 1.75TB       0 38 SSD-NVM shared - node_B_1
node_B_1:0m.i2.3L32 1.75TB       0 39 SSD-NVM shared - node_B_1
node_B_1:0n.12      1.75TB       0 12 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.13      1.75TB       0 13 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.14      1.75TB       0 14 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.15      1.75TB 0 15 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.16      1.75TB 0 16 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.17      1.75TB 0 17 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.18      1.75TB 0 18 SSD-NVM shared aggr0 node_B_1
node_B_1:0n.19      1.75TB 0 19 SSD-NVM shared - node_B_1
node_B_1:0n.24      894.0GB 0 24 SSD-NVM shared - node_A_2
node_B_1:0n.25      894.0GB 0 25 SSD-NVM shared - node_A_2
node_B_1:0n.26      894.0GB 0 26 SSD-NVM shared - node_A_2
node_B_1:0n.27      894.0GB 0 27 SSD-NVM shared - node_A_2
node_B_1:0n.28      894.0GB 0 28 SSD-NVM shared - node_A_2
node_B_1:0n.29      894.0GB 0 29 SSD-NVM shared - node_A_2
node_B_1:0n.30      894.0GB 0 30 SSD-NVM shared - node_A_2
```

```
node_B_1:0n.31      894.0GB 0 31 SSD-NVM shared - node_A_2
node_B_1:0n.36      1.75TB 0 36 SSD-NVM shared - node_A_1
node_B_1:0n.37      1.75TB 0 37 SSD-NVM shared - node_A_1
node_B_1:0n.38      1.75TB 0 38 SSD-NVM shared - node_A_1
node_B_1:0n.39      1.75TB 0 39 SSD-NVM shared - node_A_1
node_B_1:0n.40      1.75TB 0 40 SSD-NVM shared - node_A_1
node_B_1:0n.41      1.75TB 0 41 SSD-NVM shared - node_A_1
node_B_1:0n.42      1.75TB 0 42 SSD-NVM shared - node_A_1
node_B_1:0n.43      1.75TB 0 43 SSD-NVM shared - node_A_1
node_B_2:0m.i0.2L4  894.0GB 0 29 SSD-NVM shared - node_B_2
node_B_2:0m.i0.2L10 894.0GB 0 25 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L3  894.0GB 0 28 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L9  894.0GB 0 24 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L11 894.0GB 0 26 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L12 894.0GB 0 27 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L15 894.0GB 0 30 SSD-NVM shared - node_B_2
node_B_2:0m.i0.3L16 894.0GB 0 31 SSD-NVM shared - node_B_2
node_B_2:0n.0       1.75TB 0 0 SSD-NVM shared aggr0_rha12_b1_cm_02_0
node_B_2
node_B_2:0n.1 1.75TB 0 1 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.2 1.75TB 0 2 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.3 1.75TB 0 3 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.4 1.75TB 0 4 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.5 1.75TB 0 5 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.6 1.75TB 0 6 SSD-NVM shared aggr0_rha12_b1_cm_02_0 node_B_2
node_B_2:0n.7 1.75TB 0 7 SSD-NVM shared - node_B_2
64 entries were displayed.


cluster_B::>



cluster_A::> disk show
                Usable Disk Container Container
Disk Size Shelf Bay Type Type Name Owner
---------------- ---------- ----- --- ------- ---------- ---------
--------
node_A_1:0m.i1.0L2 1.75TB 0 5 SSD-NVM shared - node_B_2
node_A_1:0m.i1.0L8 1.75TB 0 3 SSD-NVM shared - node_B_2
node_A_1:0m.i1.0L18 1.75TB 0 19 SSD-NVM shared - node_B_1
node_A_1:0m.i1.0L25 1.75TB 0 12 SSD-NVM shared - node_B_1
node_A_1:0m.i1.0L27 1.75TB 0 14 SSD-NVM shared - node_B_1
node_A_1:0m.i1.2L1 1.75TB 0 4 SSD-NVM shared - node_B_2
node_A_1:0m.i1.2L6 1.75TB 0 1 SSD-NVM shared - node_B_2
node_A_1:0m.i1.2L7 1.75TB 0 2 SSD-NVM shared - node_B_2
node_A_1:0m.i1.2L14 1.75TB 0 7 SSD-NVM shared - node_B_2
node_A_1:0m.i1.2L17 1.75TB 0 18 SSD-NVM shared - node_B_1
```

```
node_A_1:0m.i1.2L22 1.75TB 0 17 SSD-NVM shared - node_B_1
node_A_1:0m.i2.1L5 1.75TB 0 0 SSD-NVM shared - node_B_2
node_A_1:0m.i2.1L13 1.75TB 0 6 SSD-NVM shared - node_B_2
node_A_1:0m.i2.1L21 1.75TB 0 16 SSD-NVM shared - node_B_1
node_A_1:0m.i2.1L26 1.75TB 0 13 SSD-NVM shared - node_B_1
node_A_1:0m.i2.1L28 1.75TB 0 15 SSD-NVM shared - node_B_1
node_A_1:0m.i2.3L19 1.75TB 0 42 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L20 1.75TB 0 43 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L23 1.75TB 0 40 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L24 1.75TB 0 41 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L29 1.75TB 0 36 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L30 1.75TB 0 37 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L31 1.75TB 0 38 SSD-NVM shared - node_A_1
node_A_1:0m.i2.3L32 1.75TB 0 39 SSD-NVM shared - node_A_1
node_A_1:0n.12 1.75TB 0 12 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.13 1.75TB 0 13 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.14 1.75TB 0 14 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.15 1.75TB 0 15 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.16 1.75TB 0 16 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.17 1.75TB 0 17 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.18 1.75TB 0 18 SSD-NVM shared aggr0 node_A_1
node_A_1:0n.19 1.75TB 0 19 SSD-NVM shared - node_A_1
node_A_1:0n.24 894.0GB 0 24 SSD-NVM shared - node_B_2
node_A_1:0n.25 894.0GB 0 25 SSD-NVM shared - node_B_2
node_A_1:0n.26 894.0GB 0 26 SSD-NVM shared - node_B_2
node_A_1:0n.27 894.0GB 0 27 SSD-NVM shared - node_B_2
node_A_1:0n.28 894.0GB 0 28 SSD-NVM shared - node_B_2
node_A_1:0n.29 894.0GB 0 29 SSD-NVM shared - node_B_2
node_A_1:0n.30 894.0GB 0 30 SSD-NVM shared - node_B_2
node_A_1:0n.31 894.0GB 0 31 SSD-NVM shared - node_B_2
node_A_1:0n.36 1.75TB 0 36 SSD-NVM shared - node_B_1
node_A_1:0n.37 1.75TB 0 37 SSD-NVM shared - node_B_1
node_A_1:0n.38 1.75TB 0 38 SSD-NVM shared - node_B_1
node_A_1:0n.39 1.75TB 0 39 SSD-NVM shared - node_B_1
node_A_1:0n.40 1.75TB 0 40 SSD-NVM shared - node_B_1
node_A_1:0n.41 1.75TB 0 41 SSD-NVM shared - node_B_1
node_A_1:0n.42 1.75TB 0 42 SSD-NVM shared - node_B_1
node_A_1:0n.43 1.75TB 0 43 SSD-NVM shared - node_B_1
node_A_2:0m.i2.3L3 894.0GB 0 28 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L4 894.0GB 0 29 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L9 894.0GB 0 24 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L10 894.0GB 0 25 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L11 894.0GB 0 26 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L12 894.0GB 0 27 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L15 894.0GB 0 30 SSD-NVM shared - node_A_2
node_A_2:0m.i2.3L16 894.0GB 0 31 SSD-NVM shared - node_A_2
```

```
node_A_2:0n.0 1.75TB 0 0 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.1 1.75TB 0 1 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.2 1.75TB 0 2 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.3 1.75TB 0 3 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.4 1.75TB 0 4 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.5 1.75TB 0 5 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.6 1.75TB 0 6 SSD-NVM shared aggr0_node_A_2_0 node_A_2
node_A_2:0n.7 1.75TB 0 7 SSD-NVM shared - node_A_2
64 entries were displayed.

cluster_A::>
```

Manually assigning drives for pool 1 (ONTAP 9.4 or later)

If the system was not preconfigured at the factory and does not meet the requirements for automatic drive assignment, you must manually assign the remote pool 1 drives.

About this task

This procedure applies to configurations running ONTAP 9.4 or later.

Details for determining whether your system requires manual disk assignment are included in Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later.

When the configuration includes only two external shelves per site, pool 1 drives for each site should be shared from the same shelf as shown in the following examples:

- node_A_1 is assigned drives in bays 0-11 on site_B-shelf_2 (remote)
- node_A_2 is assigned drives in bays 12-23 on site_B-shelf_2 (remote)

Steps

1. From each node in the MetroCluster IP configuration, assign remote drives to pool 1.

   a. Display the list of unassigned drives:

      ```
      disk show -host-adapter 0m -container-type unassigned
      ```

```
cluster_A::> disk show -host-adapter 0m -container-type unassigned
                       Usable             Disk    Container   Container
Disk                     Size Shelf Bay Type     Type         Name
Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------
6.23.0                      -    23   0 SSD     unassigned  -         -
6.23.1                      -    23   1 SSD     unassigned  -         -
.
.
.
node_A_2:0m.i1.2L51         -    21  14 SSD     unassigned  -         -
node_A_2:0m.i1.2L64         -    21  10 SSD     unassigned  -         -
.
.
.
48 entries were displayed.

cluster_A::>
```

b. Assign ownership of remote drives (0m) to pool 1 of the first node (for example, node_A_1):

   `disk assign -disk disk-id -pool 1 -owner owner-node-name`

   `disk-id` must identify a drive on a remote shelf of `owner-node-name`.

c. Confirm that the drives were assigned to pool 1:

   `disk show -host-adapter 0m -container-type unassigned`

   (i) The iSCSI connection used to access the remote drives appears as device 0m.

   The following output shows that the drives on shelf 23 were assigned because they no longer appear in the list of unassigned drives:

```
cluster_A::> disk show -host-adapter 0m -container-type unassigned
                           Usable              Disk   Container   Container
Disk                        Size Shelf Bay Type    Type        Name
Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------
node_A_2:0m.i1.2L51         -     21  14 SSD     unassigned  -          -
node_A_2:0m.i1.2L64         -     21  10 SSD     unassigned  -          -
.
.
.
node_A_2:0m.i2.1L90         -     21  19 SSD     unassigned  -          -
24 entries were displayed.

cluster_A::>
```

d. Repeat these steps to assign pool 1 drives to the second node on site A (for example, "node_A_2").

e. Repeat these steps on site B.

Manually assigning disks for pool 1 (ONTAP 9.3)

If you have at least two disk shelves for each node, you use ONTAP's auto-assignment functionality to automatically assign the remote (pool1) disks.

Before you begin

You must first assign a disk on the shelf to pool 1. ONTAP then automatically assigns the rest of the disks on the shelf to the same pool.

About this task

This procedure applies to configurations running ONTAP 9.3.

This procedure can be used only if you have at least two disk shelves for each node, which allows shelf-level auto-assignment of disks.

If you cannot use shelf-level auto-assignment, you must manually assign your remote disks so that each node has a remote pool of disks (pool 1).

The ONTAP automatic disk assignment feature assigns the disks on a shelf-by-shelf basis. For example:

• All the disks on site_B-shelf_2 are auto-assigned to pool1 of node_A_1

• All the disks on site_B-shelf_4 are auto-assigned to pool1 of node_A_2

• All the disks on site_A-shelf_2 are auto-assigned to pool1 of node_B_1

• All the disks on site_A-shelf_4 are auto-assigned to pool1 of node_B_2

You must "seed" the auto-assignment by specifying a single disk on each shelf.

Steps

1. From each node in the MetroCluster IP configuration, assign a remote disk to pool 1.

a. Display the list of unassigned disks:

```
disk show -host-adapter 0m -container-type unassigned
```

```
cluster_A::> disk show -host-adapter 0m -container-type unassigned
                         Usable              Disk     Container   Container
Disk                       Size Shelf Bay Type     Type        Name
Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------
6.23.0                        -    23   0 SSD     unassigned  -          -
6.23.1                        -    23   1 SSD     unassigned  -          -
.
.
.
node_A_2:0m.i1.2L51           -    21  14 SSD     unassigned  -          -
node_A_2:0m.i1.2L64           -    21  10 SSD     unassigned  -          -
.
.
.
48 entries were displayed.

cluster_A::>
```

b. Select a remote disk (0m) and assign ownership of the disk to pool 1 of the first node (for example, "node_A_1"):

```
disk assign -disk disk-id -pool 1 -owner owner-node-name
```

The `disk-id` must identify a disk on a remote shelf of `owner-node-name`.

The ONTAP disk auto-assignment feature assigns all disks on the remote shelf that contains the specified disk.

c. After waiting at least 60 seconds for disk auto-assignment to take place, verify that the remote disks on the shelf were auto-assigned to pool 1:

```
disk show -host-adapter 0m -container-type unassigned
```

(i) The iSCSI connection used to access the remote disks appears as device 0m.

The following output shows that the disks on shelf 23 have now been assigned and no longer appear:

```
cluster_A::> disk show -host-adapter 0m -container-type unassigned
                         Usable            Disk    Container   Container
Disk                      Size Shelf Bay Type    Type        Name
Owner
---------------- ---------- ----- --- ------- ----------- ---------
--------
node_A_2:0m.i1.2L51          -    21  14 SSD     unassigned  -          -
node_A_2:0m.i1.2L64          -    21  10 SSD     unassigned  -          -
node_A_2:0m.i1.2L72          -    21  23 SSD     unassigned  -          -
node_A_2:0m.i1.2L74          -    21   1 SSD     unassigned  -          -
node_A_2:0m.i1.2L83          -    21  22 SSD     unassigned  -          -
node_A_2:0m.i1.2L90          -    21   7 SSD     unassigned  -          -
node_A_2:0m.i1.3L52          -    21   6 SSD     unassigned  -          -
node_A_2:0m.i1.3L59          -    21  13 SSD     unassigned  -          -
node_A_2:0m.i1.3L66          -    21  17 SSD     unassigned  -          -
node_A_2:0m.i1.3L73          -    21  12 SSD     unassigned  -          -
node_A_2:0m.i1.3L80          -    21   5 SSD     unassigned  -          -
node_A_2:0m.i1.3L81          -    21   2 SSD     unassigned  -          -
node_A_2:0m.i1.3L82          -    21  16 SSD     unassigned  -          -
node_A_2:0m.i1.3L91          -    21   3 SSD     unassigned  -          -
node_A_2:0m.i2.0L49          -    21  15 SSD     unassigned  -          -
node_A_2:0m.i2.0L50          -    21   4 SSD     unassigned  -          -
node_A_2:0m.i2.1L57          -    21  18 SSD     unassigned  -          -
node_A_2:0m.i2.1L58          -    21  11 SSD     unassigned  -          -
node_A_2:0m.i2.1L59          -    21  21 SSD     unassigned  -          -
node_A_2:0m.i2.1L65          -    21  20 SSD     unassigned  -          -
node_A_2:0m.i2.1L72          -    21   9 SSD     unassigned  -          -
node_A_2:0m.i2.1L80          -    21   0 SSD     unassigned  -          -
node_A_2:0m.i2.1L88          -    21   8 SSD     unassigned  -          -
node_A_2:0m.i2.1L90          -    21  19 SSD     unassigned  -          -
24 entries were displayed.

cluster_A::>
```

    d. Repeat these steps to assign pool 1 disks to the second node on site A (for example, "node_A_2").

    e. Repeat these steps on site B.

Enabling automatic drive assignment in ONTAP 9.4

About this task

In ONTAP 9.4, if you disabled automatic drive assignment as directed previously in this procedure, you must reenable it on all nodes.

Considerations for automatic drive assignment and ADP systems in ONTAP 9.4 and later

Steps

1. Enable automatic drive assignment:

```
storage disk option modify -node node_name -autoassign on
```

You must issue this command on all nodes in the MetroCluster IP configuration.

Mirroring the root aggregates

You must mirror the root aggregates to provide data protection.

About this task

By default, the root aggregate is created as RAID-DP type aggregate. You can change the root aggregate from RAID-DP to RAID4 type aggregate. The following command modifies the root aggregate for RAID4 type aggregate:

```
storage aggregate modify -aggregate aggr_name -raidtype raid4
```

> ⓘ On non-ADP systems, the RAID type of the aggregate can be modified from the default RAID-DP to RAID4 before or after the aggregate is mirrored.

Steps

1. Mirror the root aggregate:

```
storage aggregate mirror aggr_name
```

The following command mirrors the root aggregate for "controller_A_1":

```
controller_A_1::> storage aggregate mirror aggr0_controller_A_1
```

This mirrors the aggregate, so it consists of a local plex and a remote plex located at the remote MetroCluster site.

2. Repeat the previous step for each node in the MetroCluster configuration.

Related information

Logical storage management

Creating a mirrored data aggregate on each node

You must create a mirrored data aggregate on each node in the DR group.

About this task

- You should know what drives will be used in the new aggregate.

- If you have multiple drive types in your system (heterogeneous storage), you should understand how you can ensure that the correct drive type is selected.

- Drives are owned by a specific node; when you create an aggregate, all drives in that aggregate must be owned by the same node, which becomes the home node for that aggregate.

  In systems using ADP, aggregates are created using partitions in which each drive is partitioned in to P1,

P2 and P3 partitions.

- Aggregate names should conform to the naming scheme you determined when you planned your MetroCluster configuration.

Disk and aggregate management

Steps

1. Display a list of available spares:

   `storage disk show -spare -owner node_name`

2. Create the aggregate:

   `storage aggregate create -mirror true`

   If you are logged in to the cluster on the cluster management interface, you can create an aggregate on any node in the cluster. To ensure that the aggregate is created on a specific node, use the `-node` parameter or specify drives that are owned by that node.

   You can specify the following options:

   - Aggregate's home node (that is, the node that owns the aggregate in normal operation)
   - List of specific drives that are to be added to the aggregate
   - Number of drives to include

     (i) In the minimum supported configuration, in which a limited number of drives are available, you must use the force-small-aggregate option to allow the creation of a three disk RAID-DP aggregate.

   - Checksum style to use for the aggregate
   - Type of drives to use
   - Size of drives to use
   - Drive speed to use
   - RAID type for RAID groups on the aggregate
   - Maximum number of drives that can be included in a RAID group
   - Whether drives with different RPM are allowed For more information about these options, see the storage aggregate create man page.

   The following command creates a mirrored aggregate with 10 disks:

   +

```
cluster_A::> storage aggregate create aggr1_node_A_1 -diskcount 10 -node
node_A_1 -mirror true
[Job 15] Job is queued: Create aggr1_node_A_1.
[Job 15] The job is starting.
[Job 15] Job succeeded: DONE
```

3. Verify the RAID group and drives of your new aggregate:

```
storage aggregate show-status -aggregate aggregate-name
```

Implementing the MetroCluster configuration

You must run the `metrocluster configure` command to start data protection in a MetroCluster configuration.

About this task

- There should be at least two non-root mirrored data aggregates on each cluster.

  You can verify this with the `storage aggregate show` command.

  > ⓘ   If you want to use a single mirrored data aggregate, then see Step 1 for instructions.

- The ha-config state of the controllers and chassis must be "mccip".

You issue the `metrocluster configure` command once on any of the nodes to enable the MetroCluster configuration. You do not need to issue the command on each of the sites or nodes, and it does not matter which node or site you choose to issue the command on.

The `metrocluster configure` command automatically pairs the two nodes with the lowest system IDs in each of the two clusters as disaster recovery (DR) partners. In a four-node MetroCluster configuration, there are two DR partner pairs. The second DR pair is created from the two nodes with higher system IDs.

Steps

1. Configure the MetroCluster in the following format:

If your MetroCluster configuration has…	Then do this…
Multiple data aggregates	From any node's prompt, configure MetroCluster: `metrocluster configure node-name`

A single mirrored data aggregate	a. From any node's prompt, change to the advanced privilege level:
	
	`set -privilege advanced`
	
	You need to respond with `y` when you are prompted to continue into advanced mode and you see the advanced mode prompt (*>).
	
	b. Configure the MetroCluster with the `-allow -with-one-aggregate true` parameter:
	
	`metrocluster configure -allow-with -one-aggregate true node-name`
	
	c. Return to the admin privilege level:
	
	`set -privilege admin`

(i) The best practice is to have multiple data aggregates. If the first DR group has only one aggregate and you want to add a DR group with one aggregate, you must move the metadata volume off the single data aggregate. For more information on this procedure, see Moving a metadata volume in MetroCluster configurations.

The following command enables the MetroCluster configuration on all of the nodes in the DR group that contains "controller_A_1":

```
cluster_A::*> metrocluster configure -node-name controller_A_1

[Job 121] Job succeeded: Configure is successful.
```

2. Verify the networking status on site A:

`network port show`

The following example shows the network port usage on a four-node MetroCluster configuration:

```
cluster_A::> network port show

                                                       Speed (Mbps)
Node    Port      IPspace    Broadcast Domain Link    MTU    Admin/Oper
------  --------- ---------  ---------------- -----  -------  ------------
controller_A_1
        e0a       Cluster    Cluster          up      9000   auto/1000
        e0b       Cluster    Cluster          up      9000   auto/1000
        e0c       Default    Default          up      1500   auto/1000
        e0d       Default    Default          up      1500   auto/1000
        e0e       Default    Default          up      1500   auto/1000
        e0f       Default    Default          up      1500   auto/1000
        e0g       Default    Default          up      1500   auto/1000
controller_A_2
        e0a       Cluster    Cluster          up      9000   auto/1000
        e0b       Cluster    Cluster          up      9000   auto/1000
        e0c       Default    Default          up      1500   auto/1000
        e0d       Default    Default          up      1500   auto/1000
        e0e       Default    Default          up      1500   auto/1000
        e0f       Default    Default          up      1500   auto/1000
        e0g       Default    Default          up      1500   auto/1000
14 entries were displayed.
```

3. Verify the MetroCluster configuration from both sites in the MetroCluster configuration.

   a. Verify the configuration from site A:

   metrocluster show

```
cluster_A::> metrocluster show

Configuration: IP fabric

Cluster                      Entry Name          State
-------------------------    ------------------  -----------
 Local: cluster_A            Configuration state configured
                             Mode                normal
Remote: cluster_B            Configuration state configured
                             Mode                normal
```

   b. Verify the configuration from site B:

   metrocluster show

```
cluster_B::> metrocluster show

Configuration: IP fabric

Cluster                        Entry Name          State
------------------------       ------------------  -----------
 Local: cluster_B              Configuration state configured
                               Mode                normal
Remote: cluster_A              Configuration state configured
                               Mode                normal
```

4. To avoid possible issues with nonvolatile memory mirroring, reboot each of the four nodes:

   ```
   node reboot -node node-name -inhibit-takeover true
   ```

5. Issue the `metrocluster show` command on both clusters to again verify the configuration.

Configuring the second DR group in an eight-node configuration

Repeat the previous tasks to configure the nodes in the second DR group.

Creating unmirrored data aggregates

You can optionally create unmirrored data aggregates for data that does not require the redundant mirroring provided by MetroCluster configurations.

About this task

- You should know what drives or array LUNs will be used in the new aggregate.

- If you have multiple drive types in your system (heterogeneous storage), you should understand how you can verify that the correct drive type is selected.

> (i) In MetroCluster IP configurations, remote unmirrored aggregates are not accessible after a switchover

> (i) The unmirrored aggregates must be local to the node owning them.

- Drives and array LUNs are owned by a specific node; when you create an aggregate, all drives in that aggregate must be owned by the same node, which becomes the home node for that aggregate.

- Aggregate names should conform to the naming scheme you determined when you planned your MetroCluster configuration.

- *Disks and aggregates management* contains more information about mirroring aggregates.

Steps

1. Enable unmirrored aggregate deployment:

   ```
   metrocluster modify -enable-unmirrored-aggr-deployment true
   ```

2. Verify that disk autoassignment is disabled:

```
disk option show
```

3. Install and cable the disk shelves that will contain the unmirrored aggregates.

   You can use the procedures in the Installation and Setup documentation for your platform and disk shelves.

   AFF and FAS Documentation Center

4. Manually assign all disks on the new shelf to the appropriate node:

```
disk assign -disk disk-id -owner owner-node-name
```

5. Create the aggregate:

```
storage aggregate create
```

   If you are logged in to the cluster on the cluster management interface, you can create an aggregate on any node in the cluster. To verify that the aggregate is created on a specific node, you should use the -node parameter or specify drives that are owned by that node.

   You must also ensure that you are only including drives on the unmirrored shelf to the aggregate.

   You can specify the following options:

   ◦ Aggregate's home node (that is, the node that owns the aggregate in normal operation)
   ◦ List of specific drives or array LUNs that are to be added to the aggregate
   ◦ Number of drives to include
   ◦ Checksum style to use for the aggregate
   ◦ Type of drives to use
   ◦ Size of drives to use
   ◦ Drive speed to use
   ◦ RAID type for RAID groups on the aggregate
   ◦ Maximum number of drives or array LUNs that can be included in a RAID group
   ◦ Whether drives with different RPM are allowed

     For more information about these options, see the storage aggregate create man page.

     The following command creates a unmirrored aggregate with 10 disks:

```
controller_A_1::> storage aggregate create aggr1_controller_A_1
-diskcount 10 -node controller_A_1
[Job 15] Job is queued: Create aggr1_controller_A_1.
[Job 15] The job is starting.
[Job 15] Job succeeded: DONE
```

6. Verify the RAID group and drives of your new aggregate:

```
storage aggregate show-status -aggregate aggregate-name
```

7. Disable unmirrored aggregate deployment:

```
metrocluster modify -enable-unmirrored-aggr-deployment false
```

8. Verify that disk autoassignment is enabled:

```
disk option show
```

Related information

Disk and aggregate management

Checking the MetroCluster configuration

You can check that the components and relationships in the MetroCluster configuration are working correctly.

About this task

You should do a check after initial configuration and after making any changes to the MetroCluster configuration.

You should also do a check before a negotiated (planned) switchover or a switchback operation.

If the `metrocluster check run` command is issued twice within a short time on either or both clusters, a conflict can occur and the command might not collect all data. Subsequent `metrocluster check show` commands do not show the expected output.

Steps

1. Check the configuration:

```
metrocluster check run
```

   The command runs as a background job and might not be completed immediately.

```
cluster_A::> metrocluster check run
The operation has been started and is running in the background. Wait
for
it to complete and run "metrocluster check show" to view the results. To
check the status of the running metrocluster check operation, use the
command,
"metrocluster operation history show -job-id 2245"
```

```
cluster_A::> metrocluster check show
Last Checked On: 9/13/2018 20:41:37

Component          Result
------------------ ---------
nodes              ok
lifs               ok
config-replication ok
aggregates         ok
clusters           ok
connections        ok
6 entries were displayed.
```

2. Display more detailed results from the most recent metrocluster check run command:

metrocluster check aggregate show

metrocluster check cluster show

metrocluster check config-replication show

metrocluster check lif show

metrocluster check node show

> (i) The metrocluster check show commands show the results of the most recent
> metrocluster check run command. You should always run the metrocluster
> check run command prior to using the metrocluster check show commands so that
> the information displayed is current.

The following example shows the metrocluster check aggregate show command output for a
healthy four-node MetroCluster configuration:

```
cluster_A::> metrocluster check aggregate show

Last Checked On: 8/5/2014 00:42:58

Node                    Aggregate                  Check
Result
--------------          --------------------       ---------------------
---------
controller_A_1          controller_A_1_aggr0

                                                   mirroring-status
ok

                                                   disk-pool-allocation
ok
```

```
                                                          ownership-state
ok
                    controller_A_1_aggr1
                                                 mirroring-status
ok
                                                 disk-pool-allocation
ok
                                                 ownership-state
ok
                    controller_A_1_aggr2
                                                 mirroring-status
ok
                                                 disk-pool-allocation
ok
                                                 ownership-state
ok


controller_A_2        controller_A_2_aggr0
                                                 mirroring-status
ok
                                                 disk-pool-allocation
ok
                                                 ownership-state
ok
                    controller_A_2_aggr1
                                                 mirroring-status
ok
                                                 disk-pool-allocation
ok
                                                 ownership-state
ok
                    controller_A_2_aggr2
                                                 mirroring-status
ok
                                                 disk-pool-allocation
ok
                                                 ownership-state
ok

18 entries were displayed.
```

The following example shows the metrocluster check cluster show command output for a healthy four-node MetroCluster configuration. It indicates that the clusters are ready to perform a negotiated switchover if necessary.

```
Last Checked On: 9/13/2017 20:47:04

Cluster                      Check                             Result
--------------------         ------------------------------    ---------
mccint-fas9000-0102
                             negotiated-switchover-ready       not-applicable
                             switchback-ready                  not-applicable
                             job-schedules                     ok
                             licenses                          ok
                             periodic-check-enabled            ok
mccint-fas9000-0304
                             negotiated-switchover-ready       not-applicable
                             switchback-ready                  not-applicable
                             job-schedules                     ok
                             licenses                          ok
                             periodic-check-enabled            ok
10 entries were displayed.
```

Related information

Disk and aggregate management

Network and LIF management

Completing ONTAP configuration

After configuring, enabling, and checking the MetroCluster configuration, you can proceed to complete the cluster configuration by adding additional SVMs, network interfaces and other ONTAP functionality as needed.

Verifying switchover, healing, and switchback

You should verify the switchover, healing, and switchback operations of the MetroCluster configuration.

Step

1. Use the procedures for negotiated switchover, healing, and switchback that are mentioned in the *MetroCluster Management and Disaster Recovery Guide*.

   MetroCluster management and disaster recovery

Configuring the MetroCluster Tiebreaker or ONTAP Mediator software

You can download and install on a third site either the MetroCluster Tiebreaker software, or, beginning with ONTAP 9.7, the ONTAP Mediator.

Before you begin

You must have a Linux host available that has network connectivity to both clusters in the MetroCluster configuration. The specific requirements are in the MetroCluster Tiebreaker or ONTAP Mediator documentation.

If you are connecting to an existing Tiebreaker or ONTAP Mediator instance, you need the username, password, and IP address of the Tiebreaker or Mediator service.

If you must install a new instance of the ONTAP Mediator, follow the directions to install and configure the software.

Configuring the ONTAP Mediator service for unplanned automatic switchover

If you must install a new instance of the Tiebreaker software, follow the directions to install and configure the software.

About this task

You cannot use both the MetroCluster Tiebreaker software and the ONTAP Mediator with the same MetroCluster configuration.

Considerations for using ONTAP Mediator or MetroCluster Tiebreaker

Step

1. Configure the ONTAP Mediator service or the Tiebreaker software:

   ◦ If you are using an existing instance of the ONTAP Mediator, add the ONTAP Mediator service to ONTAP:

   ```
   metrocluster configuration-settings mediator add -mediator-address ip-
   address-of-mediator-host
   ```

   ◦ If you are using the Tiebreaker software, refer to the Tiebreaker documentation.

Protecting configuration backup files

You can provide additional protection for the cluster configuration backup files by specifying a remote URL (either HTTP or FTP) where the configuration backup files will be uploaded in addition to the default locations in the local cluster.

Step

1. Set the URL of the remote destination for the configuration backup files:

   ```
   system configuration backup settings modify URL-of-destination
   ```

The Cluster Management with the CLI contains additional information under the section *Managing configuration backups*.